

COVID_finals

DS Student

5/7/2022

Finals assignment

The instructions for the assignment are to *“Import, tidy and analyze the COVID19 data set from the Johns Hopkins github site. This is the same data set I used in class. Feel free to repeat and reuse what I did if you want to. Be sure your project is reproducible and contains some visualization and analysis that is unique to your project. You may use the data to do any analysis that is of interest to you. You should include at least two visualizations and one model. Be sure to identify any bias possible in the data and in your analysis.”*

I followed the step by step process delineated in the lecture videos for week 3 to import, wrangle, analyse and visualize the data. I wrote the code in an R Markdown documents and made sure to document each step clearly so that the analysis will be reproducible. I also included the model mentioned in the lecture. Intermingled with this process I included additional analyses, visualizations, and models.

Import and read data sets

```
# import and read files

url_in <- "https://github.com/CSSEGISandData/COVID-19/raw/master/csse_covid_19_data/csse_covid_19_time_

file_names <- c("time_series_covid19_confirmed_US.csv", "time_series_covid19_confirmed_global.csv", "ti

urls <- str_c(url_in, file_names)

us_cases <- read_csv(urls[1], show_col_types = FALSE)
global_cases <- read_csv(urls[2], show_col_types = FALSE)
us_deaths <- read_csv(urls[3], show_col_types = FALSE)
global_deaths <- read_csv(urls[4], show_col_types = FALSE)
```

Explore and tidy data sets for the US

Explore us_cases

```
head(us_cases, 10)

## # A tibble: 10 x 854
##       UID iso2 iso3 code3 FIPS Admin2 Province_State Country_Region Lat
##       <dbl> <chr> <chr> <dbl> <dbl> <chr>      <chr>          <chr>      <dbl>
##  1 84001001 US    USA    840  1001 Autauga Alabama      US          32.5
##  2 84001003 US    USA    840  1003 Baldwin Alabama      US          30.7
```

```
## 3 84001005 US      USA      840 1005 Barbour Alabama      US      31.9
## 4 84001007 US      USA      840 1007 Bibb Alabama      US      33.0
## 5 84001009 US      USA      840 1009 Blount Alabama      US      34.0
## 6 84001011 US      USA      840 1011 Bullock Alabama      US      32.1
## 7 84001013 US      USA      840 1013 Butler Alabama      US      31.8
## 8 84001015 US      USA      840 1015 Calhoun Alabama      US      33.8
## 9 84001017 US      USA      840 1017 Chambers Alabama      US      32.9
## 10 84001019 US      USA      840 1019 Cherokee Alabama      US      34.2
## # ... with 845 more variables: Long_ <dbl>, Combined_Key <chr>, 1/22/20 <dbl>,
## # 1/23/20 <dbl>, 1/24/20 <dbl>, 1/25/20 <dbl>, 1/26/20 <dbl>, 1/27/20 <dbl>,
## # 1/28/20 <dbl>, 1/29/20 <dbl>, 1/30/20 <dbl>, 1/31/20 <dbl>, 2/1/20 <dbl>,
## # 2/2/20 <dbl>, 2/3/20 <dbl>, 2/4/20 <dbl>, 2/5/20 <dbl>, 2/6/20 <dbl>,
## # 2/7/20 <dbl>, 2/8/20 <dbl>, 2/9/20 <dbl>, 2/10/20 <dbl>, 2/11/20 <dbl>,
## # 2/12/20 <dbl>, 2/13/20 <dbl>, 2/14/20 <dbl>, 2/15/20 <dbl>, 2/16/20 <dbl>,
## # 2/17/20 <dbl>, 2/18/20 <dbl>, 2/19/20 <dbl>, 2/20/20 <dbl>, ...
```

Tidy us_cases

```
us_cases <- us_cases %>%
  pivot_longer(cols = -(UID:Combined_Key),
    names_to = "date",
    values_to = "cases") %>%
  select(Admin2:cases) %>%
  mutate(date = mdy(date)) %>%
  select(-c(Lat, Long_))
```

```
head(us_cases, 10)
```

```
## # A tibble: 10 x 6
##   Admin2 Province_State Country_Region Combined_Key      date      cases
##   <chr>    <chr>          <chr>          <chr>      <date>    <dbl>
## 1 Autauga Alabama        US      Autauga, Alabama, US 2020-01-22      0
## 2 Autauga Alabama        US      Autauga, Alabama, US 2020-01-23      0
## 3 Autauga Alabama        US      Autauga, Alabama, US 2020-01-24      0
## 4 Autauga Alabama        US      Autauga, Alabama, US 2020-01-25      0
## 5 Autauga Alabama        US      Autauga, Alabama, US 2020-01-26      0
## 6 Autauga Alabama        US      Autauga, Alabama, US 2020-01-27      0
## 7 Autauga Alabama        US      Autauga, Alabama, US 2020-01-28      0
## 8 Autauga Alabama        US      Autauga, Alabama, US 2020-01-29      0
## 9 Autauga Alabama        US      Autauga, Alabama, US 2020-01-30      0
## 10 Autauga Alabama        US      Autauga, Alabama, US 2020-01-31      0
```

Explore us_deaths

```
head(us_deaths)
```

```
## # A tibble: 6 x 855
##   UID iso2 iso3 code3 FIPS Admin2 Province_State Country_Region Lat
##   <dbl> <chr> <chr> <dbl> <dbl> <chr> <chr>          <chr>    <dbl>
## 1 84001001 US USA 840 1001 Autauga Alabama US 32.5
## 2 84001003 US USA 840 1003 Baldwin Alabama US 30.7
```

```
## 3 84001005 US      USA      840 1005 Barbour Alabama      US      31.9
## 4 84001007 US      USA      840 1007 Bibb    Alabama      US      33.0
## 5 84001009 US      USA      840 1009 Blount  Alabama      US      34.0
## 6 84001011 US      USA      840 1011 Bullock Alabama      US      32.1
## # ... with 846 more variables: Long_ <dbl>, Combined_Key <chr>,
## #   Population <dbl>, 1/22/20 <dbl>, 1/23/20 <dbl>, 1/24/20 <dbl>,
## #   1/25/20 <dbl>, 1/26/20 <dbl>, 1/27/20 <dbl>, 1/28/20 <dbl>, 1/29/20 <dbl>,
## #   1/30/20 <dbl>, 1/31/20 <dbl>, 2/1/20 <dbl>, 2/2/20 <dbl>, 2/3/20 <dbl>,
## #   2/4/20 <dbl>, 2/5/20 <dbl>, 2/6/20 <dbl>, 2/7/20 <dbl>, 2/8/20 <dbl>,
## #   2/9/20 <dbl>, 2/10/20 <dbl>, 2/11/20 <dbl>, 2/12/20 <dbl>, 2/13/20 <dbl>,
## #   2/14/20 <dbl>, 2/15/20 <dbl>, 2/16/20 <dbl>, 2/17/20 <dbl>, ...
```

Tidy us_deaths

```
us_deaths <- us_deaths %>%
  pivot_longer(cols = -(UID:Population),
    names_to = "date",
    values_to = "deaths") %>%
  select(Admin2:deaths) %>%
  mutate(date = mdy(date)) %>%
  select(-c(Lat, Long_))

head(us_deaths, 10)
```

```
## # A tibble: 10 x 7
##   Admin2 Province_State Country_Region Combined_Key      Population date
##   <chr>    <chr>          <chr>          <chr>          <dbl> <date>
## 1 Autauga Alabama        US      Autauga, Alabama~ 55869 2020-01-22
## 2 Autauga Alabama        US      Autauga, Alabama~ 55869 2020-01-23
## 3 Autauga Alabama        US      Autauga, Alabama~ 55869 2020-01-24
## 4 Autauga Alabama        US      Autauga, Alabama~ 55869 2020-01-25
## 5 Autauga Alabama        US      Autauga, Alabama~ 55869 2020-01-26
## 6 Autauga Alabama        US      Autauga, Alabama~ 55869 2020-01-27
## 7 Autauga Alabama        US      Autauga, Alabama~ 55869 2020-01-28
## 8 Autauga Alabama        US      Autauga, Alabama~ 55869 2020-01-29
## 9 Autauga Alabama        US      Autauga, Alabama~ 55869 2020-01-30
## 10 Autauga Alabama        US      Autauga, Alabama~ 55869 2020-01-31
## # ... with 1 more variable: deaths <dbl>
```

Join the US data sets

```
us <- us_cases %>%
  full_join(us_deaths)
```

```
## Joining, by = c("Admin2", "Province_State", "Country_Region", "Combined_Key",
## "date")
```

```
head(us, 10)
```

```
## # A tibble: 10 x 8
```

```
##   Admin2 Province_State Country_Region Combined_Key date      cases Population
##   <chr>  <chr>          <chr>          <chr>          <date>      <dbl>      <dbl>
## 1 Autau~ Alabama        US            Autauga, Al~ 2020-01-22      0      55869
## 2 Autau~ Alabama        US            Autauga, Al~ 2020-01-23      0      55869
## 3 Autau~ Alabama        US            Autauga, Al~ 2020-01-24      0      55869
## 4 Autau~ Alabama        US            Autauga, Al~ 2020-01-25      0      55869
## 5 Autau~ Alabama        US            Autauga, Al~ 2020-01-26      0      55869
## 6 Autau~ Alabama        US            Autauga, Al~ 2020-01-27      0      55869
## 7 Autau~ Alabama        US            Autauga, Al~ 2020-01-28      0      55869
## 8 Autau~ Alabama        US            Autauga, Al~ 2020-01-29      0      55869
## 9 Autau~ Alabama        US            Autauga, Al~ 2020-01-30      0      55869
## 10 Autau~ Alabama        US            Autauga, Al~ 2020-01-31      0      55869
## # ... with 1 more variable: deaths <dbl>
```

Evaluate the summary for potential problems.

```
summary(us)
```

```
##   Admin2      Province_State      Country_Region      Combined_Key
## Length:2817306 Length:2817306 Length:2817306 Length:2817306
## Class :character Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character Mode  :character
##
##
##
##      date      cases      Population      deaths
## Min.   :2020-01-22 Min.   : -3073 Min.   :      0 Min.   : -82.0
## 1st Qu.:2020-08-19 1st Qu.:   118 1st Qu.:   9917 1st Qu.:    1.0
## Median :2021-03-18 Median :   1295 Median :   24892 Median :   22.0
## Mean   :2021-03-18 Mean   :   8931 Mean   :   99604 Mean   :  140.2
## 3rd Qu.:2021-10-15 3rd Qu.:   4980 3rd Qu.:   64979 3rd Qu.:   85.0
## Max.   :2022-05-13 Max.   :2907721 Max.   :10039107 Max.   :32022.0
```

If minimum number of `cases` or `deaths` is a negative number, filter rows where either the `cases` or `deaths` variable is entered as less than 0.

```
#filter data entries where the number of cases or deaths is a negative number
us <- us %>%
  filter(cases > -1) %>%
  filter(deaths > -1)

summary(us)
```

```
##   Admin2      Province_State      Country_Region      Combined_Key
## Length:2817304 Length:2817304 Length:2817304 Length:2817304
## Class :character Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character Mode  :character
##
##
##
##      date      cases      Population      deaths
```

```
## Min. :2020-01-22 Min. : 0 Min. : 0 Min. : 0.0
## 1st Qu.:2020-08-19 1st Qu.: 118 1st Qu.: 9917 1st Qu.: 1.0
## Median :2021-03-18 Median : 1295 Median : 24909 Median : 22.0
## Mean :2021-03-17 Mean : 8931 Mean : 99604 Mean : 140.2
## 3rd Qu.:2021-10-15 3rd Qu.: 4980 3rd Qu.: 64979 3rd Qu.: 85.0
## Max. :2022-05-13 Max. :2907721 Max. :10039107 Max. :32022.0
```

Explore and tidy the global data sets

Explore global_cases

```
head(global_cases)
```

```
## # A tibble: 6 x 847
##   'Province/State' 'Country/Region' Lat Long '1/22/20' '1/23/20' '1/24/20'
##   <chr>           <chr>           <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 <NA>            Afghanistan      33.9  67.7      0      0      0
## 2 <NA>            Albania          41.2  20.2      0      0      0
## 3 <NA>            Algeria          28.0   1.66      0      0      0
## 4 <NA>            Andorra          42.5   1.52      0      0      0
## 5 <NA>            Angola          -11.2  17.9      0      0      0
## 6 <NA>            Antarctica      -71.9  23.3      0      0      0
## # ... with 840 more variables: 1/25/20 <dbl>, 1/26/20 <dbl>, 1/27/20 <dbl>,
## #   1/28/20 <dbl>, 1/29/20 <dbl>, 1/30/20 <dbl>, 1/31/20 <dbl>, 2/1/20 <dbl>,
## #   2/2/20 <dbl>, 2/3/20 <dbl>, 2/4/20 <dbl>, 2/5/20 <dbl>, 2/6/20 <dbl>,
## #   2/7/20 <dbl>, 2/8/20 <dbl>, 2/9/20 <dbl>, 2/10/20 <dbl>, 2/11/20 <dbl>,
## #   2/12/20 <dbl>, 2/13/20 <dbl>, 2/14/20 <dbl>, 2/15/20 <dbl>, 2/16/20 <dbl>,
## #   2/17/20 <dbl>, 2/18/20 <dbl>, 2/19/20 <dbl>, 2/20/20 <dbl>, 2/21/20 <dbl>,
## #   2/22/20 <dbl>, 2/23/20 <dbl>, 2/24/20 <dbl>, 2/25/20 <dbl>, ...
```

Tidy global_cases

```
global_cases <- global_cases %>%
  pivot_longer(cols = -c("Province/State", "Country/Region", "Lat", "Long"),
               names_to = "date",
               values_to = "cases") %>%
  select(-c(Lat, Long))
```

```
head(global_cases)
```

```
## # A tibble: 6 x 4
##   'Province/State' 'Country/Region' date      cases
##   <chr>           <chr>           <chr> <dbl>
## 1 <NA>            Afghanistan      1/22/20      0
## 2 <NA>            Afghanistan      1/23/20      0
## 3 <NA>            Afghanistan      1/24/20      0
## 4 <NA>            Afghanistan      1/25/20      0
## 5 <NA>            Afghanistan      1/26/20      0
## 6 <NA>            Afghanistan      1/27/20      0
```

Explore global_deaths

```
head(global_deaths)
```

```
## # A tibble: 6 x 847
##   'Province/State' 'Country/Region' Lat Long '1/22/20' '1/23/20' '1/24/20'
##   <chr>           <chr>      <dbl> <dbl> <dbl>      <dbl>      <dbl>
## 1 <NA>            Afghanistan 33.9 67.7      0        0        0
## 2 <NA>            Albania     41.2 20.2      0        0        0
## 3 <NA>            Algeria     28.0 1.66      0        0        0
## 4 <NA>            Andorra     42.5 1.52      0        0        0
## 5 <NA>            Angola     -11.2 17.9      0        0        0
## 6 <NA>            Antarctica -71.9 23.3      0        0        0
## # ... with 840 more variables: 1/25/20 <dbl>, 1/26/20 <dbl>, 1/27/20 <dbl>,
## #   1/28/20 <dbl>, 1/29/20 <dbl>, 1/30/20 <dbl>, 1/31/20 <dbl>, 2/1/20 <dbl>,
## #   2/2/20 <dbl>, 2/3/20 <dbl>, 2/4/20 <dbl>, 2/5/20 <dbl>, 2/6/20 <dbl>,
## #   2/7/20 <dbl>, 2/8/20 <dbl>, 2/9/20 <dbl>, 2/10/20 <dbl>, 2/11/20 <dbl>,
## #   2/12/20 <dbl>, 2/13/20 <dbl>, 2/14/20 <dbl>, 2/15/20 <dbl>, 2/16/20 <dbl>,
## #   2/17/20 <dbl>, 2/18/20 <dbl>, 2/19/20 <dbl>, 2/20/20 <dbl>, 2/21/20 <dbl>,
## #   2/22/20 <dbl>, 2/23/20 <dbl>, 2/24/20 <dbl>, 2/25/20 <dbl>, ...
```

Tidy global_deaths

```
global_deaths <- global_deaths %>%
  pivot_longer(cols = -c("Province/State", "Country/Region", "Lat", "Long"),
    names_to = "date",
    values_to = "deaths") %>%
  select(-c(Lat, Long))

head(global_cases, 10)
```

```
## # A tibble: 10 x 4
##   'Province/State' 'Country/Region' date      cases
##   <chr>           <chr>      <chr>      <dbl>
## 1 <NA>            Afghanistan 1/22/20      0
## 2 <NA>            Afghanistan 1/23/20      0
## 3 <NA>            Afghanistan 1/24/20      0
## 4 <NA>            Afghanistan 1/25/20      0
## 5 <NA>            Afghanistan 1/26/20      0
## 6 <NA>            Afghanistan 1/27/20      0
## 7 <NA>            Afghanistan 1/28/20      0
## 8 <NA>            Afghanistan 1/29/20      0
## 9 <NA>            Afghanistan 1/30/20      0
## 10 <NA>           Afghanistan 1/31/20      0
```

Join the global data sets

full_join the two global data sets and rename two of the columns. mutate date to a date object.

```
global <- global_cases %>%
  full_join(global_deaths) %>%
  rename(Province_State = "Province/State",
    Country_Region = "Country/Region") %>%
  mutate(date = mdy(date))
```

```
## Joining, by = c("Province/State", "Country/Region", "date")
```

```
head(global, 10)
```

```
## # A tibble: 10 x 5
##   Province_State Country_Region date      cases deaths
##   <chr>          <chr>      <date>    <dbl>  <dbl>
## 1 <NA>          Afghanistan 2020-01-22      0      0
## 2 <NA>          Afghanistan 2020-01-23      0      0
## 3 <NA>          Afghanistan 2020-01-24      0      0
## 4 <NA>          Afghanistan 2020-01-25      0      0
## 5 <NA>          Afghanistan 2020-01-26      0      0
## 6 <NA>          Afghanistan 2020-01-27      0      0
## 7 <NA>          Afghanistan 2020-01-28      0      0
## 8 <NA>          Afghanistan 2020-01-29      0      0
## 9 <NA>          Afghanistan 2020-01-30      0      0
## 10 <NA>         Afghanistan 2020-01-31      0      0
```

Evaluate the summary for potential problems.

```
summary(global)
```

```
## Province_State      Country_Region      date      cases
## Length:239412      Length:239412      Min.   :2020-01-22      Min.   :      0
## Class :character    Class :character    1st Qu.:2020-08-19      1st Qu.:     284
## Mode  :character    Mode  :character    Median :2021-03-18      Median :     6340
##                      Mean   :2021-03-18      Mean   :    554666
##                      3rd Qu.:2021-10-15      3rd Qu.:   121880
##                      Max.   :2022-05-13      Max.   :  82421624
##
## deaths
## Min.   :      0
## 1st Qu.:      2
## Median :     82
## Mean   :   10151
## 3rd Qu.:   1900
## Max.   :  999518
```

The minimum number of cases is zero. `filter` out rows with no cases.

```
global <- global %>%
  filter(cases > 0)
```

```
summary(global)
```

```
## Province_State      Country_Region      date      cases
## Length:220713      Length:220713      Min.   :2020-01-22      Min.   :      1
## Class :character    Class :character    1st Qu.:2020-09-24      1st Qu.:     693
## Mode  :character    Mode  :character    Median :2021-04-14      Median :    10127
##                      Mean   :2021-04-10      Mean   :   601658
##                      3rd Qu.:2021-10-29      3rd Qu.:  149729
##                      Max.   :2022-05-13      Max.   :  82421624
```

```
##      deaths
## Min.   :    0
## 1st Qu.:    5
## Median :   125
## Mean   : 11011
## 3rd Qu.: 2399
## Max.   :999518
```

Now the minimum number of cases in any row is one.

Check to ensure the Max in each of the `cases` and `deaths` columns is not a typographical error.

```
global %>% filter(cases > 80000000)
```

```
## # A tibble: 47 x 5
##   Province_State Country_Region date       cases deaths
##   <chr>          <chr>      <date>    <dbl>  <dbl>
## 1 <NA>          US        2022-03-28 80001286 978260
## 2 <NA>          US        2022-03-29 80025464 979243
## 3 <NA>          US        2022-03-30 80064646 980411
## 4 <NA>          US        2022-03-31 80110284 980927
## 5 <NA>          US        2022-04-01 80142499 981756
## 6 <NA>          US        2022-04-02 80154308 981912
## 7 <NA>          US        2022-04-03 80159063 981970
## 8 <NA>          US        2022-04-04 80187971 982557
## 9 <NA>          US        2022-04-05 80213572 983070
## 10 <NA>         US        2022-04-06 80254519 984195
## # ... with 37 more rows
```

Check the `deaths` column.

```
global %>% filter(deaths > 990000)
```

```
## # A tibble: 24 x 5
##   Province_State Country_Region date       cases deaths
##   <chr>          <chr>      <date>    <dbl>  <dbl>
## 1 <NA>          US        2022-04-20 80804471 990330
## 2 <NA>          US        2022-04-21 80864015 990713
## 3 <NA>          US        2022-04-22 80954071 991269
## 4 <NA>          US        2022-04-23 80973785 991320
## 5 <NA>          US        2022-04-24 80987251 991349
## 6 <NA>          US        2022-04-25 81057431 991634
## 7 <NA>          US        2022-04-26 81102716 991958
## 8 <NA>          US        2022-04-27 81192697 992759
## 9 <NA>          US        2022-04-28 81267145 993126
## 10 <NA>         US        2022-04-29 81328234 993553
## # ... with 14 more rows
```

Add population data to global data set and a `Combined_Key` variable

`Combined_Key` variable combines `Province_State` and `Country_Region` into one variable.


```
global <- global %>%
  unite("Combined_Key",
        c(Province_State, Country_Region),
        # combines the column names with a space
        sep = " ",
        na.rm = TRUE,
        remove = FALSE)
```

Add Population from a csv file at the same Johns Hopkins website

```
uid_lookup_url <- "https://raw.githubusercontent.com/CSSEGISandData/COVID-19/master/csse_covid_19_data/"
```

Read the csv file in and look at the column names

```
uid <- read_csv(uid_lookup_url, show_col_types = FALSE)

colnames(uid)
```

```
## [1] "UID"          "iso2"          "iso3"          "code3"
## [5] "FIPS"         "Admin2"        "Province_State" "Country_Region"
## [9] "Lat"          "Long_"         "Combined_Key"  "Population"
```

select desired columns

```
uid <- uid %>%
  select(-c(Lat, Long_, Combined_Key, code3, iso2, iso3, Admin2))

head(uid, 10)
```

```
## # A tibble: 10 x 5
##   UID FIPS Province_State Country_Region Population
##   <dbl> <chr> <chr>          <chr>          <dbl>
## 1     4 <NA> <NA>          Afghanistan    38928341
## 2     8 <NA> <NA>          Albania        2877800
## 3    10 <NA> <NA>          Antarctica         NA
## 4    12 <NA> <NA>          Algeria        43851043
## 5    20 <NA> <NA>          Andorra         77265
## 6    24 <NA> <NA>          Angola        32866268
## 7    28 <NA> <NA>          Antigua and Barbuda 97928
## 8    32 <NA> <NA>          Argentina    45195777
## 9    51 <NA> <NA>          Armenia       2963234
## 10   40 <NA> <NA>          Austria       9006400
```

left_join the global population data set to the global COVID data set. Reorder the new data set columns (use select).

```
global <- global %>%
  left_join(uid, by = c("Province_State", "Country_Region")) %>%
  select(-c(UID, FIPS)) %>%
  select(Province_State, Country_Region, date, cases, deaths, Population, Combined_Key)

head(global, 10)
```

```
## # A tibble: 10 x 7
##   Province_State Country_Region date       cases deaths Population Combined_Key
##   <chr>          <chr>      <date>    <dbl>  <dbl>      <dbl> <chr>
## 1 <NA>          Afghanistan 2020-02-24     5      0    38928341 Afghanistan
## 2 <NA>          Afghanistan 2020-02-25     5      0    38928341 Afghanistan
## 3 <NA>          Afghanistan 2020-02-26     5      0    38928341 Afghanistan
## 4 <NA>          Afghanistan 2020-02-27     5      0    38928341 Afghanistan
## 5 <NA>          Afghanistan 2020-02-28     5      0    38928341 Afghanistan
## 6 <NA>          Afghanistan 2020-02-29     5      0    38928341 Afghanistan
## 7 <NA>          Afghanistan 2020-03-01     5      0    38928341 Afghanistan
## 8 <NA>          Afghanistan 2020-03-02     5      0    38928341 Afghanistan
## 9 <NA>          Afghanistan 2020-03-03     5      0    38928341 Afghanistan
## 10 <NA>         Afghanistan 2020-03-04     5      0    38928341 Afghanistan
```

Although the Province_State column may be empty for many countries, it is reported for some. We can explore this by printing out the unique Province_State values in the global data set.

```
head(unique(global$Province_State), 15)
```

```
## [1] NA "Australian Capital Territory"
## [3] "New South Wales" "Northern Territory"
## [5] "Queensland" "South Australia"
## [7] "Tasmania" "Victoria"
## [9] "Western Australia" "Alberta"
## [11] "British Columbia" "Diamond Princess"
## [13] "Grand Princess" "Manitoba"
## [15] "New Brunswick"
```

Analyse the global data

To get the number of daily cases and deaths by country we need to `group_by` Province_State, Country_Region, and date and then `sum` the cases, deaths, and Population for each country. Calculate death rate as deaths per million and add as a column (using `mutate`). `select` the column names to include. `ungroup` the data set.

```
global_summary <- global %>%
  group_by(Province_State, Country_Region, date) %>%
  summarize(cases = sum(cases), deaths = sum(deaths),
            Population = sum(Population)) %>%
  mutate(deaths_per_mill = deaths * 1000000 / Population) %>%
  select(Country_Region, date, cases, deaths, deaths_per_mill, Population) %>%
  ungroup() %>%
  select(-Province_State)
```

```
## 'summarise()' has grouped output by 'Province_State', 'Country_Region'. You can override using the '
## Adding missing grouping variables: 'Province_State'
```

```
tail(global_summary, 10)
```

```
## # A tibble: 10 x 6
##   Country_Region date       cases deaths deaths_per_mill Population
```

```
##      <chr>          <date>      <dbl> <dbl>          <dbl>      <dbl>
##  1 Zimbabwe      2022-05-04 248050  5471          368.    14862927
##  2 Zimbabwe      2022-05-05 248050  5471          368.    14862927
##  3 Zimbabwe      2022-05-06 248214  5473          368.    14862927
##  4 Zimbabwe      2022-05-07 248214  5473          368.    14862927
##  5 Zimbabwe      2022-05-08 248352  5476          368.    14862927
##  6 Zimbabwe      2022-05-09 248536  5479          369.    14862927
##  7 Zimbabwe      2022-05-10 248642  5481          369.    14862927
##  8 Zimbabwe      2022-05-11 248642  5481          369.    14862927
##  9 Zimbabwe      2022-05-12 248943  5481          369.    14862927
## 10 Zimbabwe      2022-05-13 249131  5482          369.    14862927
```

Look at the data for France

```
France <- global_summary %>%
  filter(cases > 0 & Country_Region == 'France')
tail(France, 10)
```

```
## # A tibble: 10 x 6
##   Country_Region date      cases deaths deaths_per_mill Population
##   <chr>          <date>      <dbl> <dbl>          <dbl>      <dbl>
##  1 France      2022-05-04 28026502 143176          2194.    65249843
##  2 France      2022-05-05 28070727 143298          2196.    65249843
##  3 France      2022-05-06 28108929 143408          2198.    65249843
##  4 France      2022-05-07 28146887 143470          2199.    65249843
##  5 France      2022-05-08 28176211 143524          2200.    65249843
##  6 France      2022-05-09 28181004 143655          2202.    65249843
##  7 France      2022-05-10 28229588 143772          2203.    65249843
##  8 France      2022-05-11 28269887 143868          2205.    65249843
##  9 France      2022-05-12 28305934 143952          2206.    65249843
## 10 France      2022-05-13 28334484 144048          2208.    65249843
```

Analyze the US data

Analyze data for a state - group by state Look at the column names for the US data set again

```
colnames(us)
```

```
## [1] "Admin2"          "Province_State" "Country_Region" "Combined_Key"
## [5] "date"           "cases"          "Population"      "deaths"
```

filter the data for WA state and look at the data set

```
us_WA <- us %>%
  filter(Province_State == 'Washington')
head(us_WA, 10)
```

```
## # A tibble: 10 x 8
##   Admin2 Province_State Country_Region Combined_Key date      cases Population
```

```
##      <chr> <chr>          <chr>          <chr>          <date>      <dbl>      <dbl>
## 1 Adams Washington US Adams, Wash~ 2020-01-22 0 19983
## 2 Adams Washington US Adams, Wash~ 2020-01-23 0 19983
## 3 Adams Washington US Adams, Wash~ 2020-01-24 0 19983
## 4 Adams Washington US Adams, Wash~ 2020-01-25 0 19983
## 5 Adams Washington US Adams, Wash~ 2020-01-26 0 19983
## 6 Adams Washington US Adams, Wash~ 2020-01-27 0 19983
## 7 Adams Washington US Adams, Wash~ 2020-01-28 0 19983
## 8 Adams Washington US Adams, Wash~ 2020-01-29 0 19983
## 9 Adams Washington US Adams, Wash~ 2020-01-30 0 19983
## 10 Adams Washington US Adams, Wash~ 2020-01-31 0 19983
## # ... with 1 more variable: deaths <dbl>
```

Each state contains multiple counties. To get the number of daily cases and deaths by state we need to group_by Province_State, Country_Region, and date and then sum the cases, deaths, and Population for the counties comprising each state. Calculate death rate as deaths per million and add as a column (using mutate). select the column names to include. ungroup the data set.

```
us_by_state <- us %>%
  group_by(Province_State, Country_Region, date) %>%
  summarize(cases = sum(cases), deaths = sum(deaths),
            Population = sum(Population)) %>%
  mutate(deaths_per_mill = deaths * 1000000 / Population) %>%
  select(Province_State, Country_Region, date, cases, deaths, deaths_per_mill, Population) %>%
  ungroup()
```

'summarise()' has grouped output by 'Province_State', 'Country_Region'. You can
override using the '.groups' argument.

```
tail(us_by_state, 10)
```

```
## # A tibble: 10 x 7
##   Province_State Country_Region date      cases deaths deaths_per_mill
##   <chr>          <chr>      <date>      <dbl> <dbl>      <dbl>
## 1 Wyoming      US      2022-05-04 156745 1814      3134.
## 2 Wyoming      US      2022-05-05 156745 1814      3134.
## 3 Wyoming      US      2022-05-06 156745 1814      3134.
## 4 Wyoming      US      2022-05-07 156745 1814      3134.
## 5 Wyoming      US      2022-05-08 156745 1814      3134.
## 6 Wyoming      US      2022-05-09 156745 1814      3134.
## 7 Wyoming      US      2022-05-10 157031 1817      3139.
## 8 Wyoming      US      2022-05-11 157031 1817      3139.
## 9 Wyoming      US      2022-05-12 157031 1817      3139.
## 10 Wyoming     US      2022-05-13 157031 1817      3139.
## # ... with 1 more variable: Population <dbl>
```

Look at the data for WA state in the last 10 days.

```
us_WA <- us_by_state %>%
  filter(Province_State == 'Washington')

tail(us_WA, 10)
```

```
## # A tibble: 10 x 7
##   Province_State Country_Region date       cases deaths deaths_per_mill
##   <chr>          <chr>      <date>    <dbl> <dbl>      <dbl>
## 1 Washington    US        2022-05-04 1508065 12716      1670.
## 2 Washington    US        2022-05-05 1508065 12716      1670.
## 3 Washington    US        2022-05-06 1508065 12716      1670.
## 4 Washington    US        2022-05-07 1508065 12716      1670.
## 5 Washington    US        2022-05-08 1508065 12716      1670.
## 6 Washington    US        2022-05-09 1519327 12742      1673.
## 7 Washington    US        2022-05-10 1519327 12742      1673.
## 8 Washington    US        2022-05-11 1524078 12770      1677.
## 9 Washington    US        2022-05-12 1524078 12770      1677.
## 10 Washington   US        2022-05-13 1530430 12791      1680.
## # ... with 1 more variable: Population <dbl>
```

Analyze data for the US - group all states together Analyze the daily cases and deaths in the entirety of the US. Look at the data for the US in the last 10 days.

```
us_total <- us %>%
  group_by(Country_Region, date) %>%
  summarize(cases = sum(cases), deaths = sum(deaths),
            Population = sum(Population)) %>%
  mutate(deaths_per_mill = deaths * 1000000 / Population) %>%
  select(Country_Region, date, cases, deaths, deaths_per_mill, Population) %>%
  ungroup()
```

```
## 'summarise()' has grouped output by 'Country_Region'. You can override using the
## '.groups' argument.
```

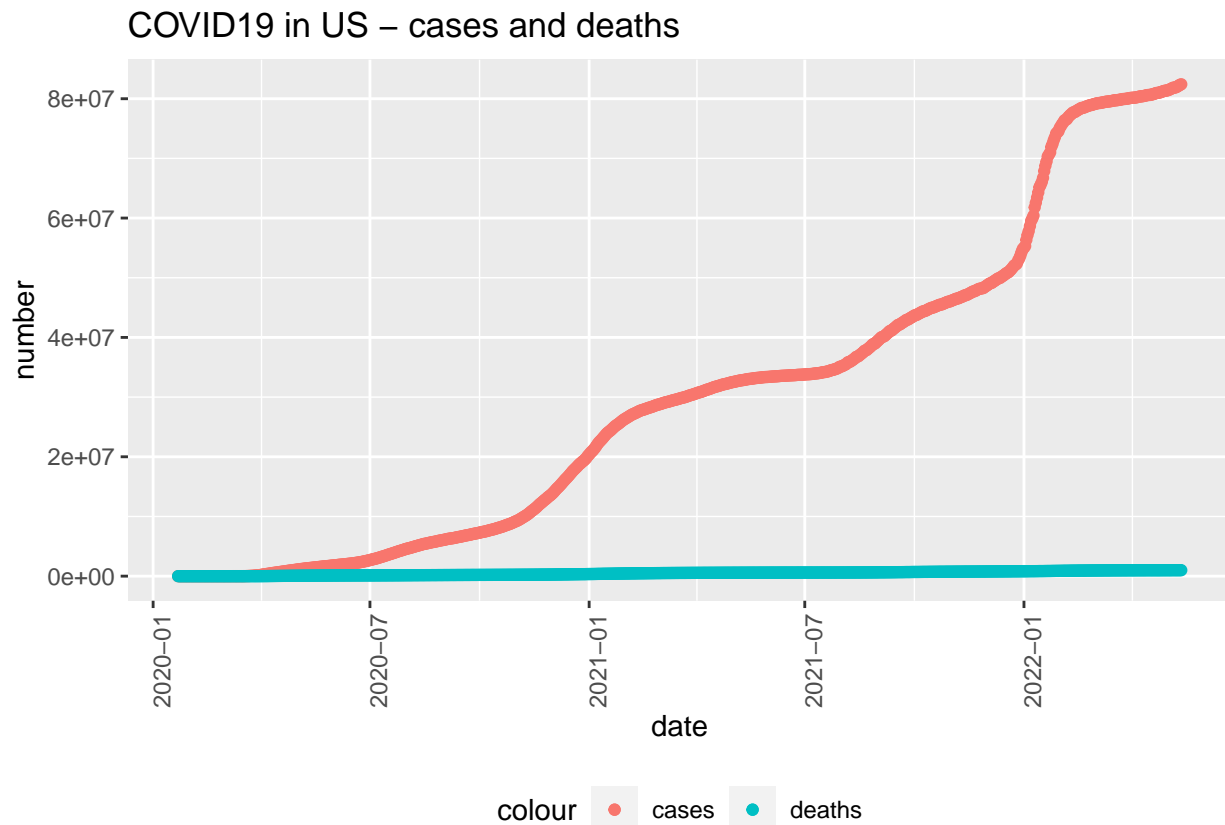
```
tail(us_total, 10)
```

```
## # A tibble: 10 x 6
##   Country_Region date       cases deaths deaths_per_mill Population
##   <chr>          <date>    <dbl> <dbl>      <dbl>      <dbl>
## 1 US            2022-05-04 81620724 996656      2994. 332875137
## 2 US            2022-05-05 81710925 997046      2995. 332875137
## 3 US            2022-05-06 81834681 997400      2996. 332875137
## 4 US            2022-05-07 81858498 997503      2997. 332875137
## 5 US            2022-05-08 81863479 997526      2997. 332875137
## 6 US            2022-05-09 81973661 997740      2997. 332875137
## 7 US            2022-05-10 82059839 998048      2998. 332875137
## 8 US            2022-05-11 82223174 998997      3001. 332875137
## 9 US            2022-05-12 82325687 999125      3002. 332875137
## 10 US           2022-05-13 82421624 999518      3003. 332875137
```

Visualize the US data

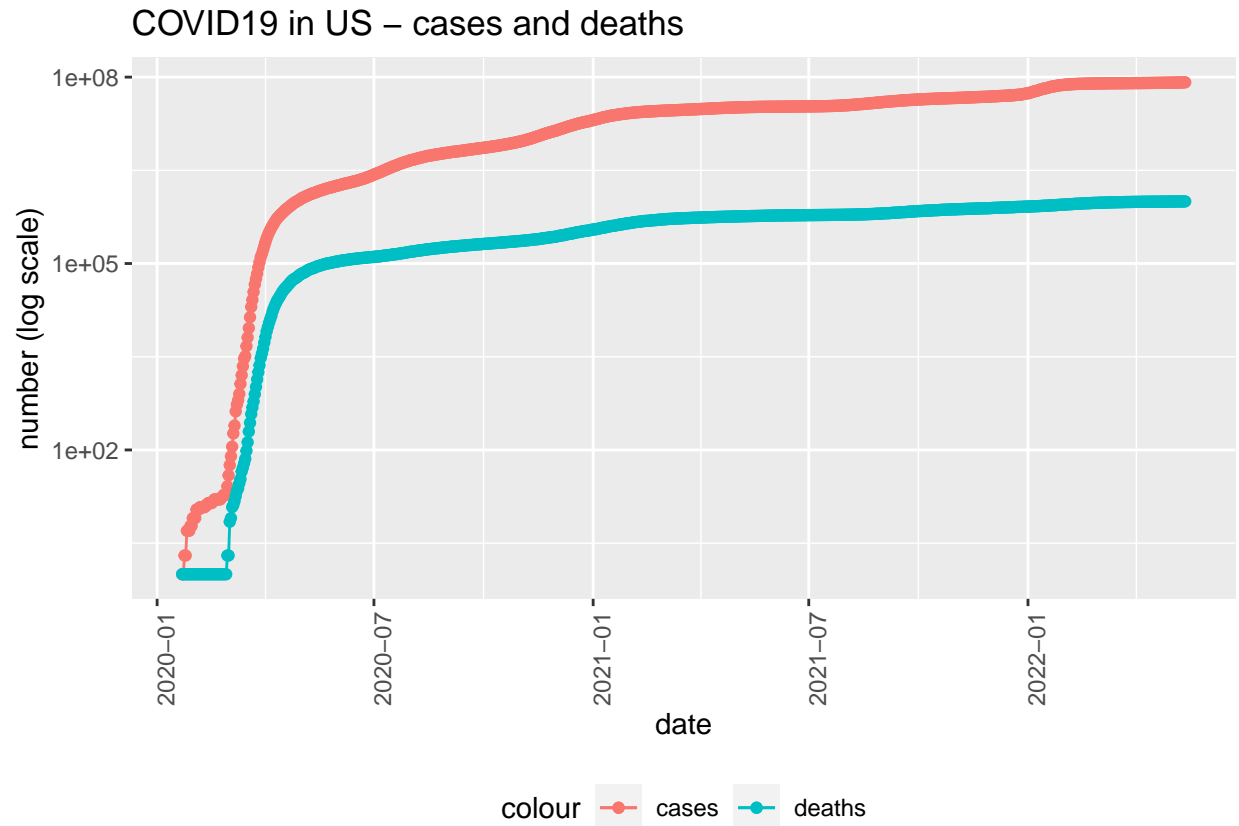
filter the `us_total` to only include dates with cases. Plot the `cases` as `geom_point`. Plot the `deaths` as `geom_point`.

```
us_total %>%
  filter(cases > 0) %>%
  ggplot(aes(x = date, y = cases)) +
  geom_point(aes(color = "cases")) +
  geom_point(aes(y = deaths, color = "deaths")) +
  theme(legend.position="bottom",
        axis.text.x = element_text(angle = 90)) +
  labs(title = "COVID19 in US - cases and deaths", y = "number")
```



Optimize the plot by: - Plotting the cases and deaths as `geom_line` as well. - Changing the y-axis to a logarithmic scale (`scale_y_log`).

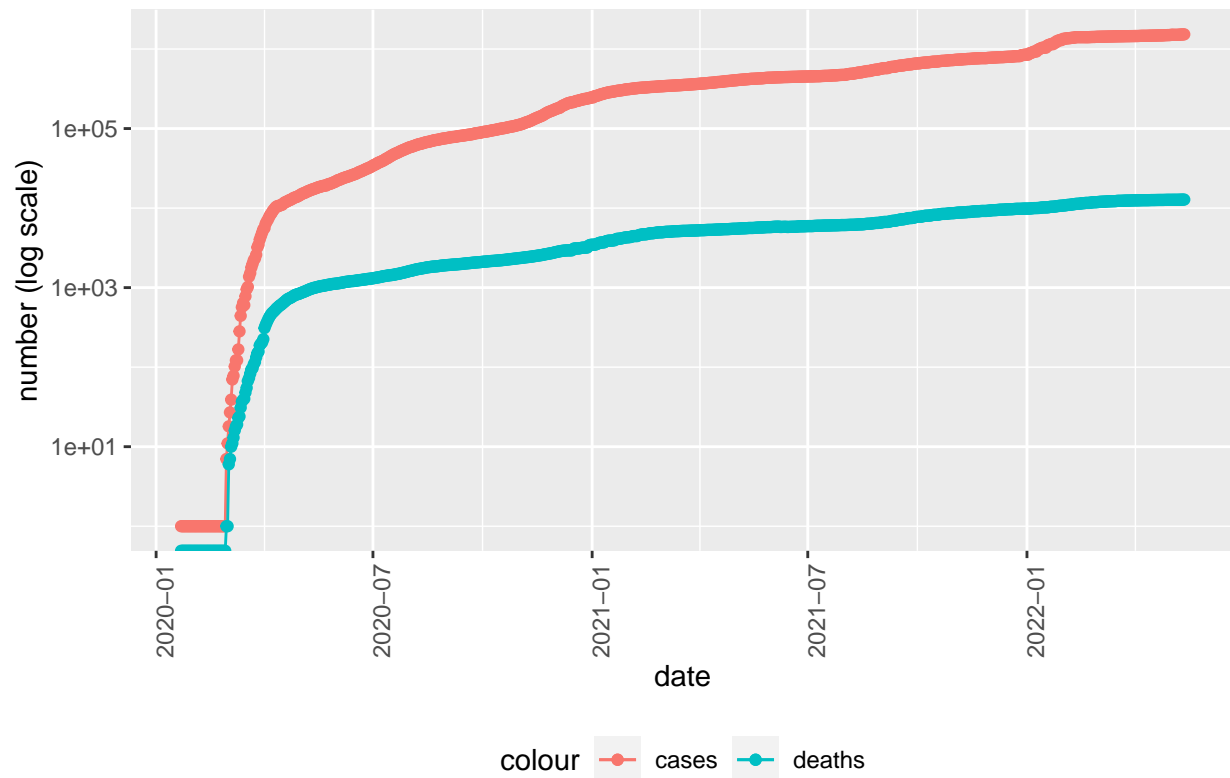
```
us_total %>%
  filter(cases > 0) %>%
  ggplot(aes(x = date, y = cases)) +
  geom_line(aes(color = "cases")) +
  geom_point(aes(color = "cases")) +
  geom_line(aes(y = deaths, color = "deaths")) +
  geom_point(aes(y = deaths, color = "deaths")) +
  scale_y_log10() +
  theme(legend.position="bottom",
        axis.text.x = element_text(angle = 90)) +
  labs(title = "COVID19 in US - cases and deaths", y = "number (log scale)")
```



Visualize the data for *WA state*.

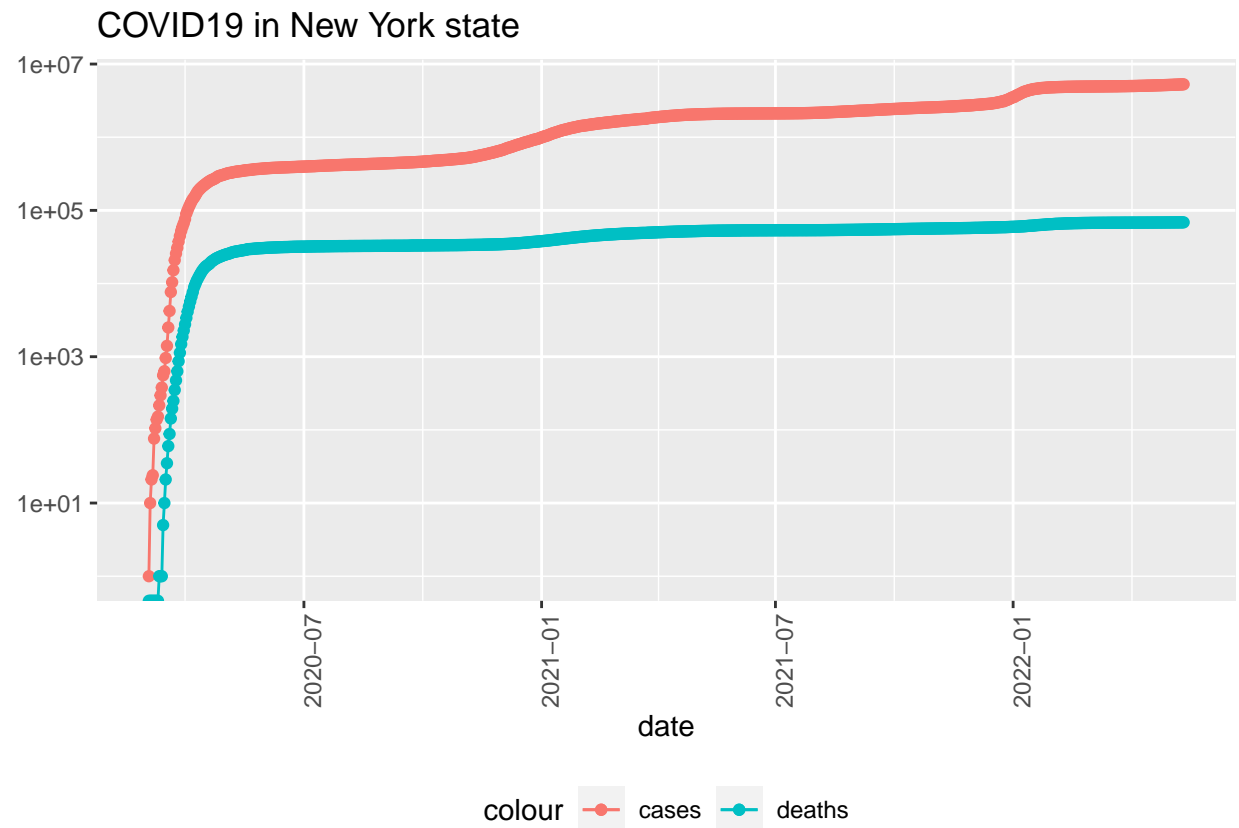
```
us_by_state %>%
  filter(cases > 0 & Province_State == 'Washington') %>%
  ggplot(aes(x = date, y = cases)) +
  geom_line(aes(color = "cases")) +
  geom_point(aes(color = "cases")) +
  geom_line(aes(y = deaths, color = "deaths")) +
  geom_point(aes(y = deaths, color = "deaths")) +
  scale_y_log10() +
  theme(legend.position="bottom",
        axis.text.x = element_text(angle = 90)) +
  labs(title = "COVID19 in WA state - cases and deaths", y = "number (log scale)")
```

COVID19 in WA state – cases and deaths



Visualize the data for *New York state*.

```
us_by_state %>%
  filter(cases > 0 & Province_State == 'New York') %>%
  ggplot(aes(x = date, y = cases)) +
  geom_line(aes(color = "cases")) +
  geom_point(aes(color = "cases")) +
  geom_line(aes(y = deaths, color = "deaths")) +
  geom_point(aes(y = deaths, color = "deaths")) +
  scale_y_log10() +
  theme(legend.position="bottom",
        axis.text.x = element_text(angle = 90)) +
  labs(title = "COVID19 in New York state", y = NULL)
```

The visualizations suggest that the number of cases deaths have leveled off. Is this true?

Analyzing data

Explore the `us_total` data set: - current date - max total deaths - max deaths per million - max popultion

```
max(us_total$date)
```

```
## [1] "2022-05-13"
```

```
max(us_total$deaths)
```

```
## [1] 999518
```

```
max(us_total$deaths_per_mill)
```

```
## [1] 3002.681
```

```
max(us_total$Population)
```

```
## [1] 332875137
```

Explore the us_by_state data set for the same data: - current date - max total deaths - max deaths per million - max population

```
max(us_by_state$date)
```

```
## [1] "2022-05-13"
```

```
max(us_by_state$deaths)
```

```
## [1] 90782
```

```
max(us_by_state$deaths_per_mill)
```

```
## [1] NaN
```

```
max(us_by_state$Population)
```

```
## [1] 39512223
```

Transform data - add new variables

Add new variables: - us_by_states add: new_cases and new_deaths - us_total add: new_cases and new_deaths

```
us_by_state <- us_by_state %>%
  mutate(new_cases = cases - lag(cases),
         new_deaths = deaths - lag(deaths))

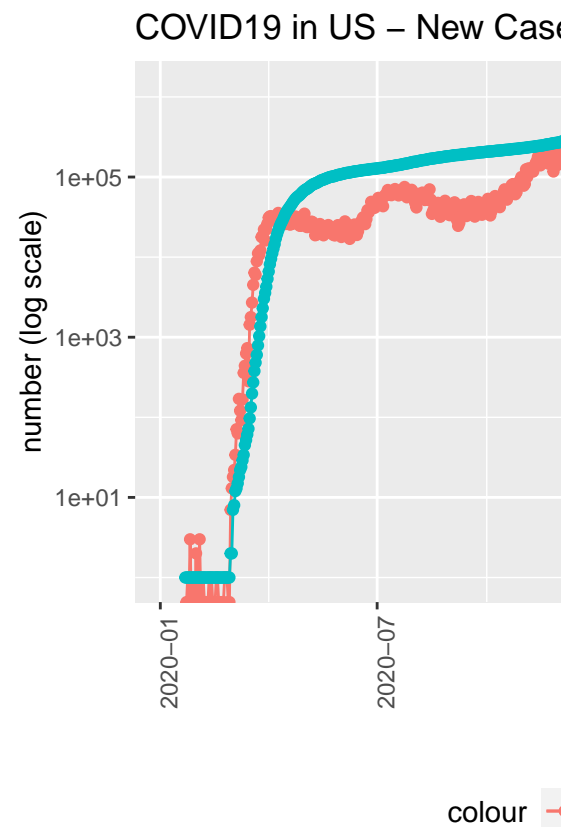
us_total <- us_total %>%
  mutate(new_cases = cases - lag(cases),
         new_deaths = deaths - lag(deaths))

tail(us_total %>%
  select(new_cases, new_deaths, everything()), 10)
```

```
## # A tibble: 10 x 8
##   new_cases new_deaths Country_Region date      cases deaths deaths_per_mill
##   <dbl>      <dbl> <chr>      <date>      <dbl> <dbl>      <dbl>
## 1  114107      1959 US        2022-05-04  8.16e7 996656    2994.
## 2   90201       390 US        2022-05-05  8.17e7 997046    2995.
## 3  123756       354 US        2022-05-06  8.18e7 997400    2996.
## 4   23817       103 US        2022-05-07  8.19e7 997503    2997.
## 5    4981        23 US        2022-05-08  8.19e7 997526    2997.
## 6  110182       214 US        2022-05-09  8.20e7 997740    2997.
## 7   86178       308 US        2022-05-10  8.21e7 998048    2998.
## 8  163335       949 US        2022-05-11  8.22e7 998997    3001.
## 9   102513      128 US        2022-05-12  8.23e7 999125    3002.
## 10    95937      393 US        2022-05-13  8.24e7 999518    3003.
## # ... with 1 more variable: Population <dbl>
```

Visualize transformed data

```
us_total %>%
  ggplot(aes(x = date, y = new_cases)) +
  geom_line(aes(color = "new_cases")) +
  geom_point(aes(color = "new_cases")) +
  geom_line(aes(y = deaths, color = "new_deaths")) +
  geom_point(aes(y = deaths, color = "new_deaths")) +
  scale_y_log10() +
  theme(legend.position="bottom",
        axis.text.x = element_text(angle = 90)) +
  labs(title = "COVID19 in US - New Cases & New Deaths", y = "number (log scale)")
```

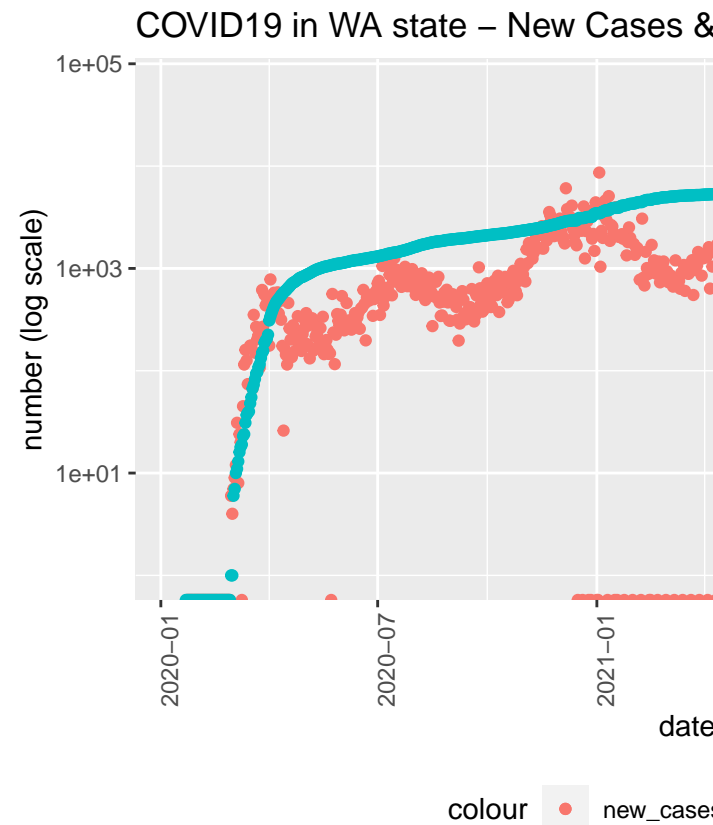


Visualize the new_cases and new_deaths in the us_total data set

Observation: The plot demonstrates that the daily have flattened and there are fluctuations in the number of daily new_cases. Last peak of daily new_cases - the largest peak so far - was in January 2022.

```
us_by_state %>%
  filter(cases > 0 & Province_State == 'Washington') %>%
  ggplot(aes(x = date, y = new_cases)) +
  #geom_line(aes(color = "new_cases")) +
  geom_point(aes(color = "new_cases")) +
```

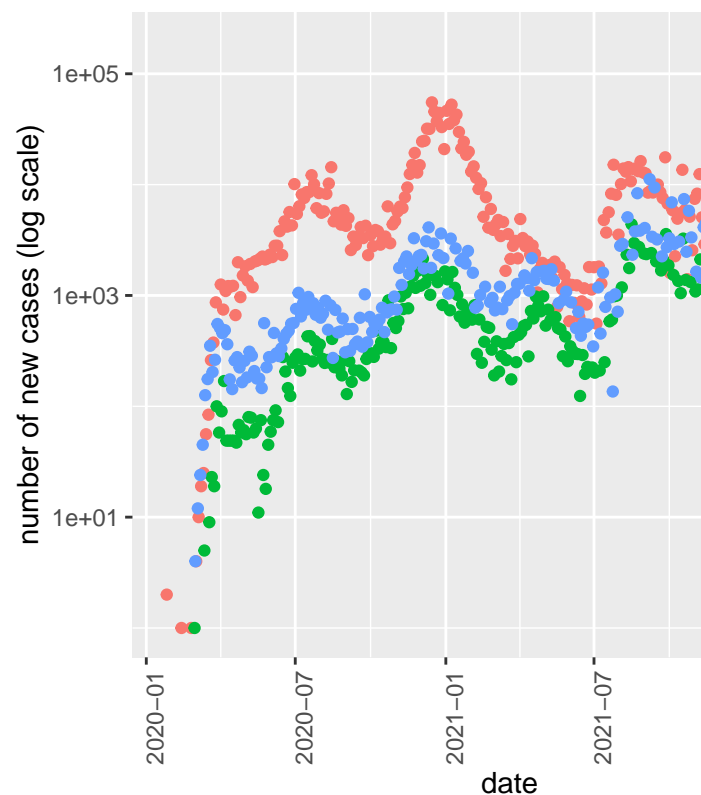
```
#geom_line(aes(y = deaths, color = "new_deaths")) +
geom_point(aes(y = deaths, color = "new_deaths")) +
scale_y_log10() +
theme(legend.position="bottom",
      axis.text.x = element_text(angle = 90)) +
labs(title = "COVID19 in WA state - New Cases & New Deaths", y = "number (log scale)")
```



Visualize the new_cases and new_deaths in WA state

```
us_by_state %>%
  filter(new_cases > 0 & Province_State == c("Washington", "California", "Oregon") ) %>%
  ggplot(aes(x = date, y = new_cases, group = Province_State, color = factor(Province_State))) +
  geom_point() +
  scale_y_log10() +
  theme(legend.position="right",
        axis.text.x = element_text(angle = 90)) +
  labs(title = "COVID19 in WA, CA, and OR states - New Cases", y = "number of new cases (log scale)")
```

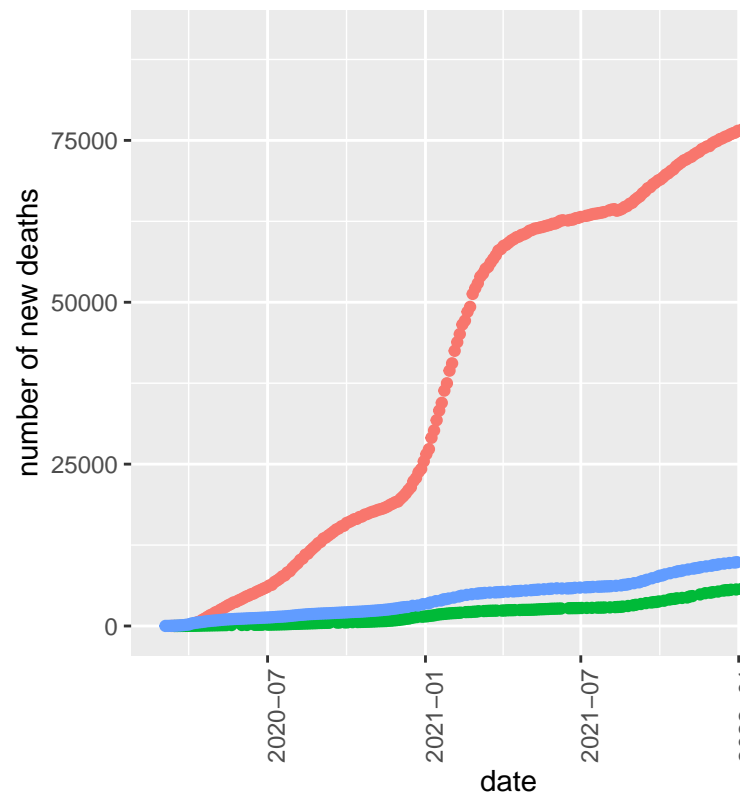
COVID19 in WA, CA, and OR states



Visualize the trend of new cases in three western states

```
# number of new deaths is numeric (i.e., not a log scale)
us_by_state %>%
  filter(new_deaths > 0 & Province_State == c("Washington", "California", "Oregon") ) %>%
  ggplot(aes(x = date, y = deaths, group = Province_State, color = factor(Province_State))) +
  geom_point() +
  #scale_y_log10() +
  theme(legend.position="right",
        axis.text.x = element_text(angle = 90)) +
  labs(title = "COVID19 in WA, CA, and OR states - New Deaths", y = "number of new deaths")
```

COVID19 in WA, CA, and OR states – M



Visualize the trend of deaths in three western states

Analyse data: What are the best and worst states?

Transform data to a new table which summarizes the `deaths`, `cases`, `cases_per_thou`, and `deaths_per_thou` in each state. `arrange` the rows in ascending number of `deaths` in each state.

```
us_state_totals <- us_by_state %>%
  group_by(Province_State) %>%
  summarize(deaths = max(deaths), cases = max(cases),
            population = max(Population),
            cases_per_thou = 1000 * cases / population,
            deaths_per_thou = 1000 * deaths / population) %>%
  filter(cases > 0, population > 0) %>%
  arrange(deaths)
```

us_state_totals

```
## # A tibble: 56 x 6
##   Province_State    deaths    cases population cases_per_thou deaths_per_thou
##   <chr>          <dbl>   <dbl>    <dbl>         <dbl>         <dbl>
## 1 American Samoa      30    5999    55641          108.           0.539
## 2 Northern Mariana Isl~  34   11305    55144          205.           0.617
## 3 Virgin Islands     113   18037   107268          168.           1.05
## 4 Guam              361  48496   164229          295.           2.20
## 5 Vermont           648 128471   623989          206.           1.04
## 6 Alaska          1268 256089   740995          346.           1.71
```

```
## 7 District of Columbia      1340 144675      705749      205.      1.90
## 8 Hawaii                    1434 258422      1415872      183.      1.01
## 9 Wyoming                   1817 157031      578759       271.      3.14
## 10 North Dakota             2272 242462      762062      318.      2.98
## # ... with 46 more rows
```

The 10 states with the **lowest death rates per 1000 population**.

```
us_state_totals %>%
  slice_min(deaths_per_thou, n = 10) %>%
  select(Province_State, deaths_per_thou, cases_per_thou, everything())
```

```
## # A tibble: 10 x 6
##   Province_State deaths_per_thou cases_per_thou deaths cases population
##   <chr>          <dbl>          <dbl> <dbl> <dbl> <dbl>
## 1 American Samoa      0.539          108.    30 6.00e3  55641
## 2 Northern Mariana Isl~ 0.617          205.    34 1.13e4  55144
## 3 Hawaii              1.01          183.   1434 2.58e5 1415872
## 4 Vermont             1.04          206.    648 1.28e5  623989
## 5 Virgin Islands      1.05          168.   113 1.80e4  107268
## 6 Puerto Rico         1.13          158.  4250 5.94e5 3754939
## 7 Utah                1.49          293.  4761 9.39e5 3205958
## 8 Washington          1.68          201. 12791 1.53e6 7614893
## 9 Alaska              1.71          346.  1268 2.56e5  740995
## 10 Maine              1.74          189.  2335 2.54e5 1344212
```

The 10 states with the **highest death rates per 1000 population**.

```
us_state_totals %>%
  slice_max(deaths_per_thou, n = 10) %>%
  select(Province_State, deaths_per_thou, cases_per_thou, everything())
```

```
## # A tibble: 10 x 6
##   Province_State deaths_per_thou cases_per_thou deaths cases population
##   <chr>          <dbl>          <dbl> <dbl> <dbl> <dbl>
## 1 Mississippi       4.19          269. 12457 801527 2976149
## 2 Arizona            4.15          279. 30230 2030925 7278717
## 3 Oklahoma           4.04          264. 15996 1044179 3956971
## 4 Alabama            4.00          266. 19628 1304710 4903185
## 5 West Virginia      3.85          282.  6893 505528 1792147
## 6 Tennessee          3.85          298. 26265 2036315 6829174
## 7 Arkansas           3.78          278. 11415 838251 3017804
## 8 New Jersey         3.78          260. 33537 2313062 8882190
## 9 Louisiana          3.72          254. 17295 1178806 4648794
## 10 New Mexico        3.63          251.  7607 526137 2096829
```

The 10 states with the **highest case rates per 1000 population**.

```
us_state_totals %>%
  slice_max(cases_per_thou, n = 10) %>%
  select(Province_State, cases_per_thou, deaths_per_thou, everything())
```

```
## # A tibble: 10 x 6
##   Province_State cases_per_thou deaths_per_thou deaths    cases population
##   <chr>          <dbl>          <dbl> <dbl>    <dbl>    <dbl>
## 1 Rhode Island      360.            3.35  3552  381271  1059361
## 2 Alaska            346.            1.71  1268  256089   740995
## 3 North Dakota      318.            2.98  2272  242462   762062
## 4 Kentucky          299.            3.54 15797 1336858  4467673
## 5 Tennessee         298.            3.85 26265 2036315  6829174
## 6 Guam              295.            2.20   361   48496   164229
## 7 Utah              293.            1.49  4761  938864  3205958
## 8 South Carolina    288.            3.47 17869 1481975  5148714
## 9 West Virginia     282.            3.85  6893  505528  1792147
## 10 Wisconsin        282.            2.49 14502 1639365  5822434
```

```
us_state_totals
```

```
## # A tibble: 56 x 6
##   Province_State deaths cases population cases_per_thou deaths_per_thou
##   <chr>          <dbl> <dbl>    <dbl>          <dbl>          <dbl>
## 1 American Samoa      30   5999   55641         108.           0.539
## 2 Northern Mariana Isl~  34  11305   55144         205.           0.617
## 3 Virgin Islands      113 18037  107268         168.           1.05
## 4 Guam                361 48496  164229         295.           2.20
## 5 Vermont             648 128471  623989         206.           1.04
## 6 Alaska            1268 256089  740995         346.           1.71
## 7 District of Columbia 1340 144675  705749         205.           1.90
## 8 Hawaii             1434 258422 1415872         183.           1.01
## 9 Wyoming            1817 157031  578759         271.           3.14
## 10 North Dakota       2272 242462  762062         318.           2.98
## # ... with 46 more rows
```

My Additional Analyses and Visualizations In these analyses I tried to filter and limit the number of states to approximately 10 so that the visualizations are not too crowded. I looked at the states with the highest death rates per 1000 and highest case rates per thousand.

Visualize the **total deaths per state** for those states that have a total death of *greater than 33000*. First, filter the data and analyze it in a table; **arrange** in ascending order of number of deaths.

```
deaths_states_1 <- us_state_totals %>%
  filter(deaths > 33000) %>%
  arrange(deaths) %>%
  select(Province_State, deaths)

deaths_states_1
```

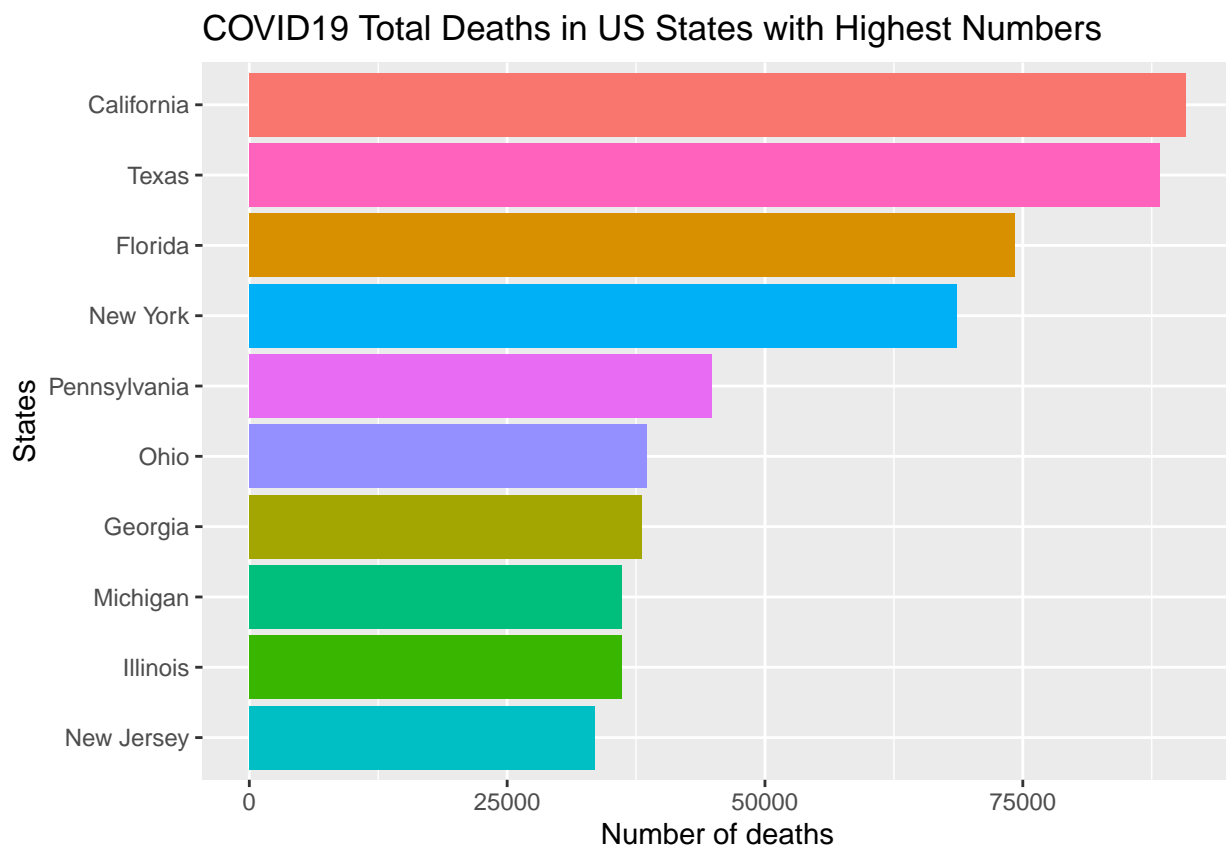
```
## # A tibble: 10 x 2
##   Province_State deaths
##   <chr>          <dbl>
## 1 New Jersey     33537
## 2 Illinois       36138
## 3 Michigan       36140
## 4 Georgia        38086
## 5 Ohio           38550
```



```
## 6 Pennsylvania 44814
## 7 New York 68603
## 8 Florida 74178
## 9 Texas 88240
## 10 California 90782
```

Visualize the above table as a bar graph (with `geom_col()`). Use `fct_reorder()` from `forcats` package to reorder the States in descending order of number of total deaths.

```
deaths_states_1 %>%
  ggplot(aes(x = fct_reorder(Province_State, deaths), y = deaths, group = Province_State, fill = Province_State)) +
  geom_col(show.legend = FALSE) +
  coord_flip() +
  labs(x = "States", y = "Number of deaths", title = "COVID19 Total Deaths in US States with Highest Numbers")
```



```
deaths_states_1
```

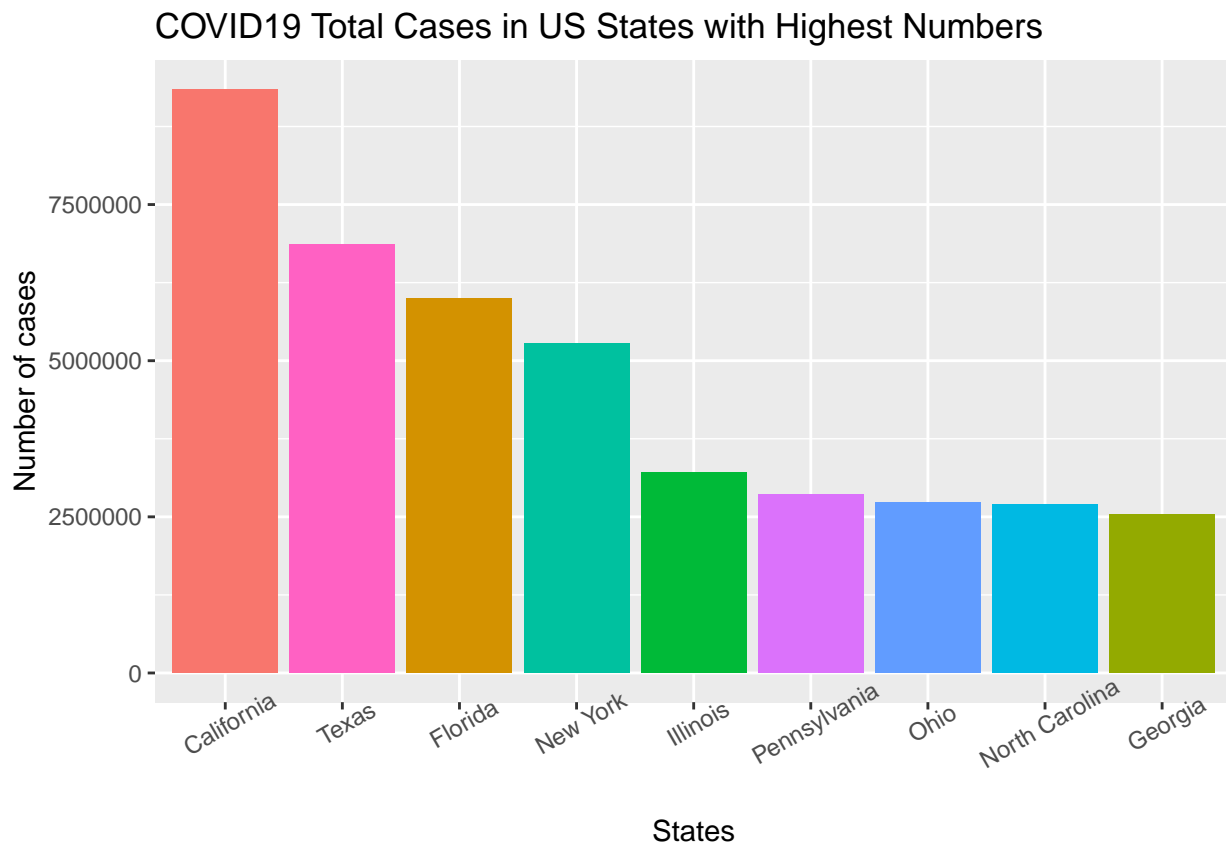
```
## # A tibble: 10 x 2
##   Province_State deaths
##   <chr>          <dbl>
## 1 New Jersey    33537
## 2 Illinois     36138
## 3 Michigan     36140
## 4 Georgia      38086
## 5 Ohio         38550
```

```
## 6 Pennsylvania 44814
## 7 New York 68603
## 8 Florida 74178
## 9 Texas 88240
## 10 California 90782
```

Reverse ordering , i.e., ascending, is achieved with `fct_rev()`.

```
cases_states <- us_state_totals %>%
  filter(cases > 2500000) %>%
  ggplot(aes(x = fct_rev(fct_reorder(Province_State, cases)), y = cases, group = Province_State, fill = Province_State)) +
  geom_col(show.legend = FALSE) +
  theme(legend.position="bottom",
        axis.text.x = element_text(angle = 30)) +
  labs(x = "States", y = "Number of cases", title = "COVID19 Total Cases in US States with Highest Numbers")

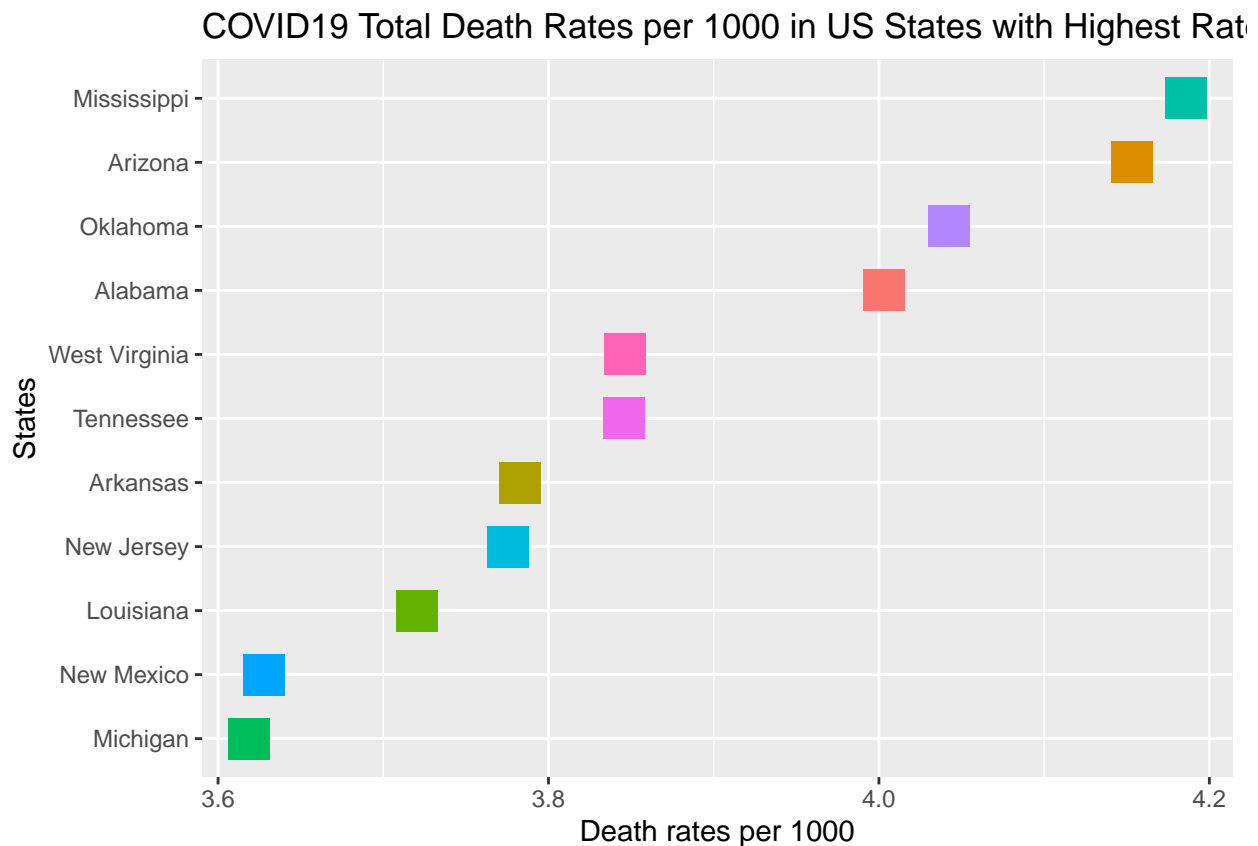
cases_states
```



Visualize the states with the **highest death rates per 1000**.

```
# need to improve the labeling
deaths_states_2 <- us_state_totals %>%
  filter(deaths_per_thou > 3.6) %>%
  ggplot(aes(x = fct_reorder(Province_State, deaths_per_thou), y = deaths_per_thou, group = Province_State, fill = Province_State)) +
  geom_point(show.legend = FALSE, size = 7, shape = 15) +
```

```
coord_flip() +
theme(legend.position="bottom",
      axis.text.x = element_text(angle = 0)) +
labs(title = "COVID19 Total Death Rates per 1000 in US States with Highest Rates", x = "States", y =
deaths_states_2
```



Question: What is the trend of `new_deaths` in the states which have had the highest death rates in the last 2 years?

At what time points during the pandemic have they had the highest rate of `new_deaths`?

Has the number of `new_deaths` fluctuated with adopted preventative measures and policies and vaccinations in these states? To answer this question we need more data regarding the type and timeline of the policies, vaccination administration as well as the introduction of new COVID strains in the specific locale.

```
# filtered the deaths_per_thou to greater than 3.9, so that I would have less than 5 states to visualize

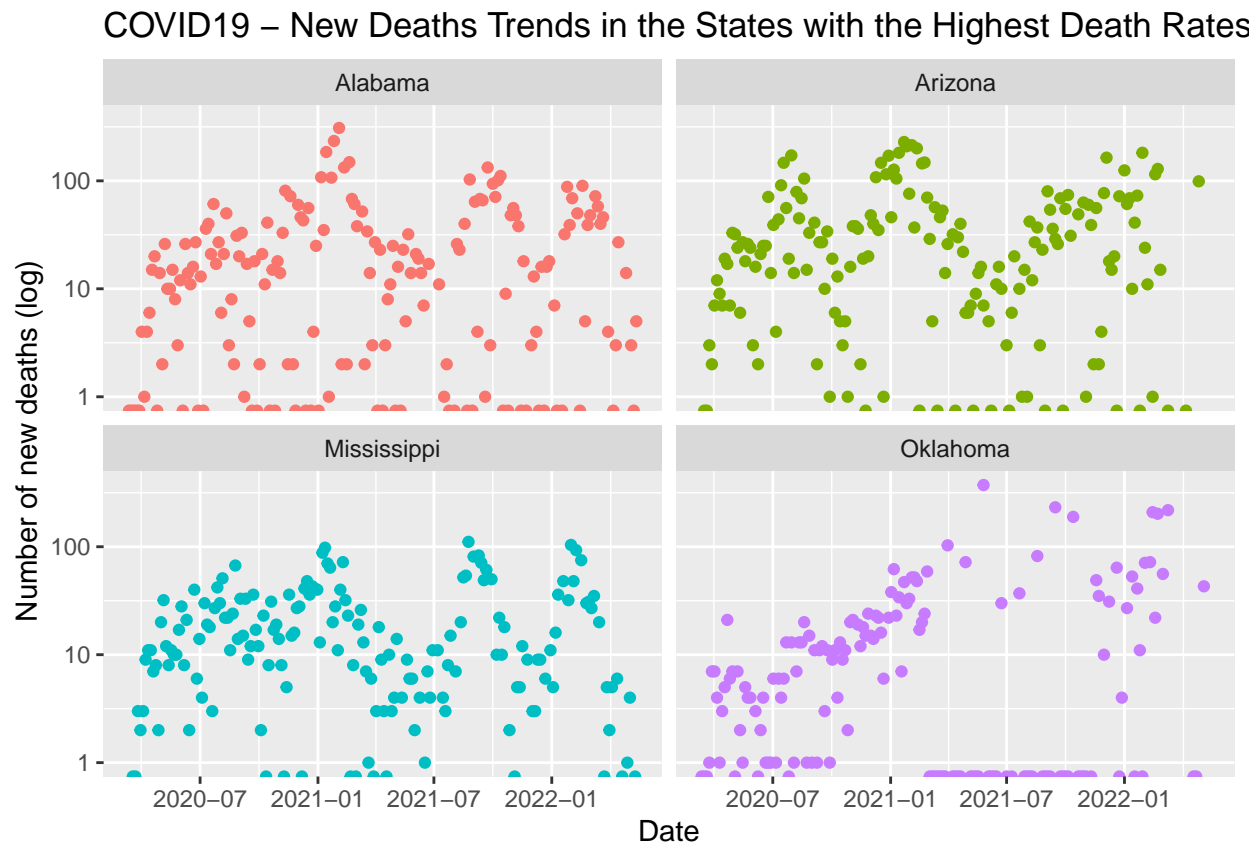
states_highest_death_rate <- us_state_totals %>%
  filter(deaths_per_thou > 3.9) %>%
  select(Province_State)

# use deframe() to change the tibble to a vector
states_highest = deframe(states_highest_death_rate)
states_highest
```

```
## [1] "Mississippi" "Oklahoma" "Alabama" "Arizona"
```

Plot the new_deaths in these states.

```
us_by_state %>%
  filter(new_cases > 0 & Province_State == states_highest) %>%
  ggplot(aes(x = date, y = new_deaths, group = Province_State, color = factor(Province_State))) +
  geom_point(show.legend = FALSE) +
  scale_y_log10() +
  # used facet_wrap to graph the data for each state individually
  facet_wrap(~ Province_State, ncol = 2) +
  labs(title = "COVID19 - New Deaths Trends in the States with the Highest Death Rates", y = "Number of
```



Question: What is the trend of new_deaths in the states which have had the lowest death rates in the last 2 years?

How do these compare to the states which have had the highest death rates?

```
# filtered the deaths_per_thou to greater than 3.9, so that I would have less than 5 states to visualize

states_lowest_death_rate <- us_state_totals %>%
  filter(deaths_per_thou < 1.5) %>%
  select(Province_State)

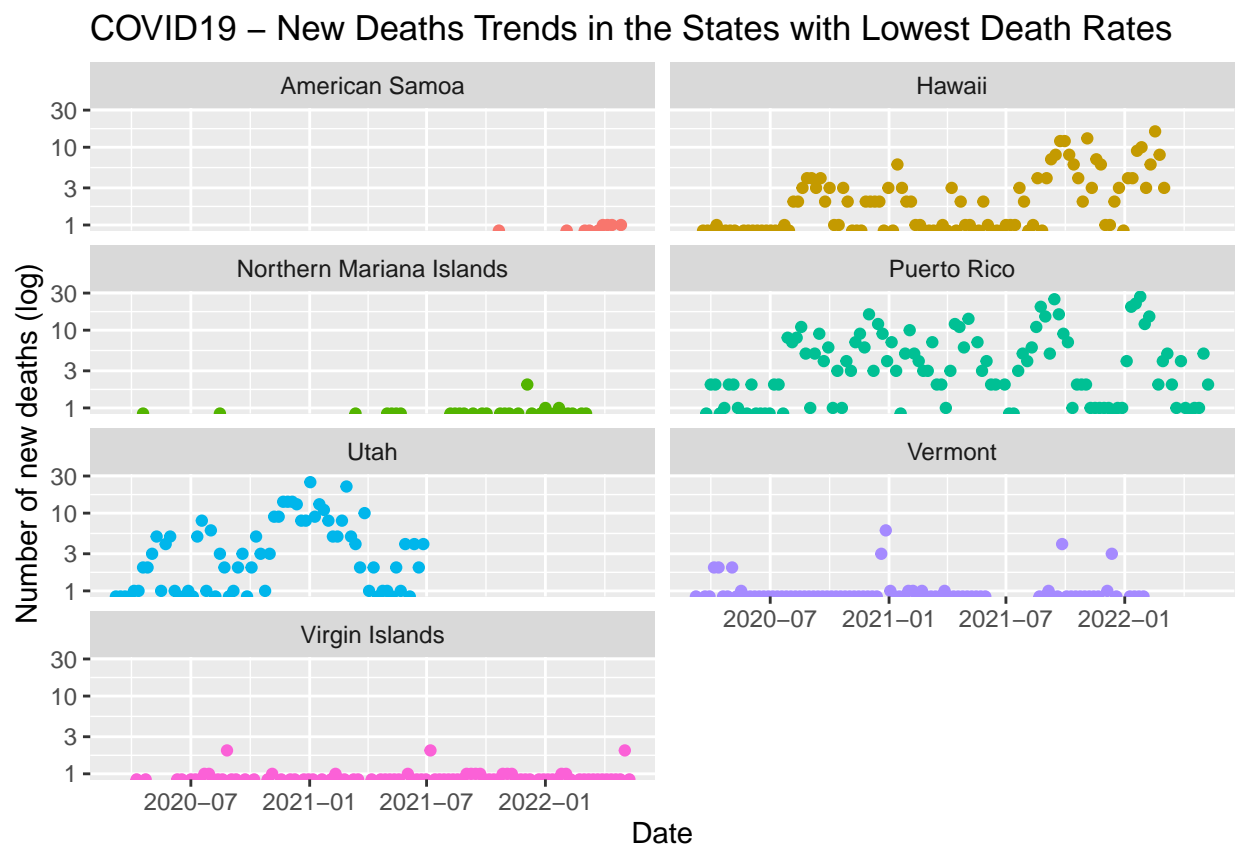
# use deframe() to change the tibble to a vector
states_lowest = deframe(states_lowest_death_rate)
states_lowest
```

```
## [1] "American Samoa" "Northern Mariana Islands"
```

```
## [3] "Virgin Islands"      "Vermont"
## [5] "Hawaii"              "Puerto Rico"
## [7] "Utah"
```

Plot the `new_deaths` in these states.

```
us_by_state %>%
  filter(new_cases > 0 & Province_State == states_lowest) %>%
  ggplot(aes(x = date, y = new_deaths, group = Province_State, color = factor(Province_State))) +
  geom_point(show.legend = FALSE) +
  scale_y_log10() +
  facet_wrap(~ Province_State, ncol = 2) +
  labs(title = "COVID19 - New Deaths Trends in the States with Lowest Death Rates", y = "Number of new c
```



Observation: Except for the state with very low numbers of new deaths, in both the states with the highest and lowest new deaths we can discern a cyclical pattern of high and lows that may correspond to the emergence of the COVID19 virus strains or specific national events (such as travel and gatherings during the Christmas and New Year holidays).

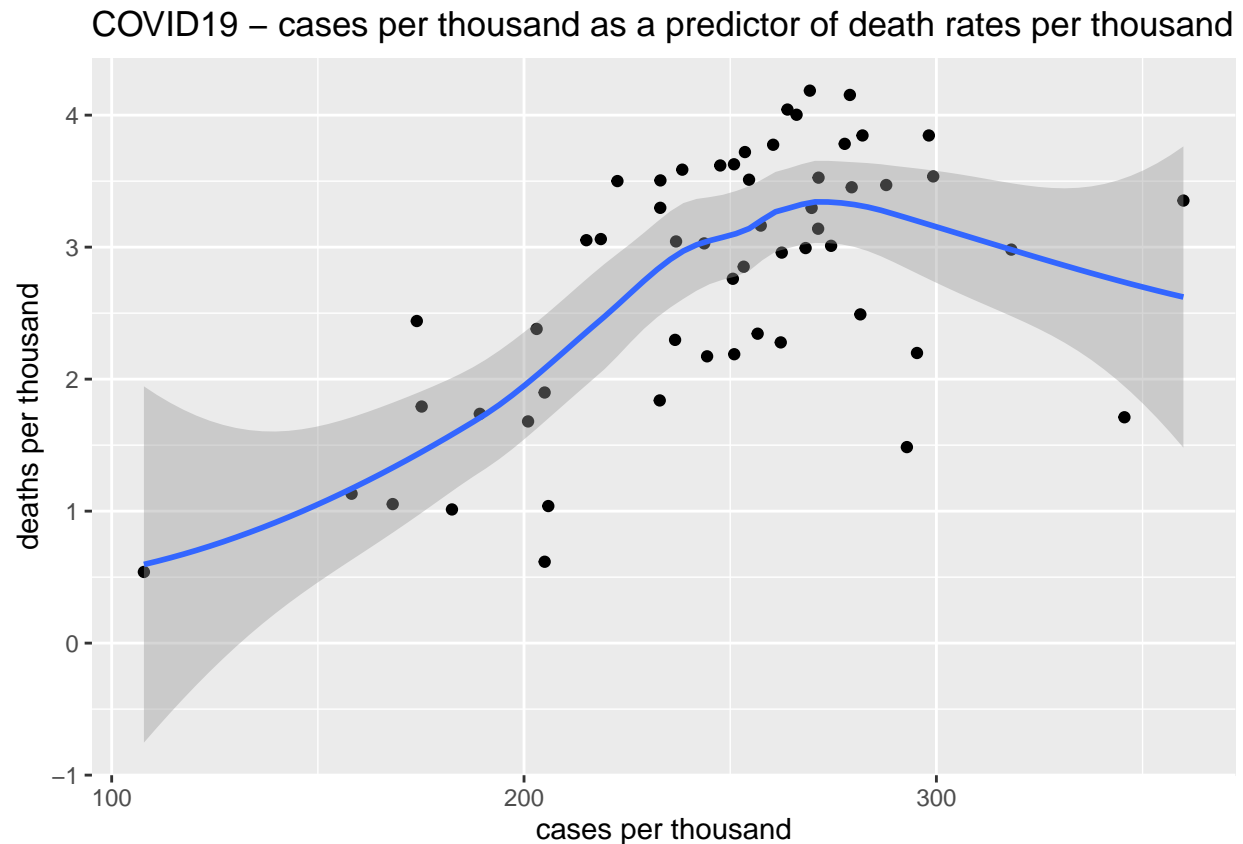
Modeling the Data - part 1

As enumerated in the week 3 lecture on 'Modeling Data' some variables that can be added and considered for to the model include population density, climate of the area, political affiliation, extent of the lock down, etc.

Model 1: `cases_per_thou` as a predictor for `deaths_per_thou` (a linear model)

Plot the two variables as a scatter plot to see their relationship.

```
us_state_totals %>% ggplot(aes(cases_per_thou,deaths_per_thou)) +  
  geom_point() +  
  geom_smooth() +  
  labs(title = "COVID19 - cases per thousand as a predictor of death rates per thousand", x = "cases per  
  
## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```



The linear model ...

```
mod_1 <- lm(deaths_per_thou ~ cases_per_thou, data = us_state_totals)  
  
summary(mod_1)  
  
##  
## Call:  
## lm(formula = deaths_per_thou ~ cases_per_thou, data = us_state_totals)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -2.2440 -0.5617  0.1102  0.6547  1.1506   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept)    -0.213801    0.594120   -0.360    0.72
## cases_per_thou  0.012063    0.002365    5.102 4.47e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7968 on 54 degrees of freedom
## Multiple R-squared:  0.3252, Adjusted R-squared:  0.3127
## F-statistic: 26.03 on 1 and 54 DF,  p-value: 4.469e-06
```

We can calculate `slice_min` and `slice_max` of the variable `cases_per_thou` to determine the range of values in `cases_per_thou`.

```
min_cases <- us_state_totals %>%
  slice_min(cases_per_thou) %>%
  select(cases_per_thou)

min_cases = as.integer(min_cases)
min_cases
```

```
## [1] 107
```

```
max_cases <- us_state_totals %>%
  slice_max(cases_per_thou) %>%
  select(cases_per_thou)

max_cases = as.integer(max_cases)
max_cases
```

```
## [1] 359
```

The `cases_per_thou` variable therefore ranges from 107 to 359.

`mutate` a new variable for the predicted death rate per thousand (`pred`) and `arrange` the tibble' in ascending order for the value of `pred`.

Look at the 10 highest predicted death rates using `tail`.

```
us_state_w_pred <- us_state_totals %>%
  mutate(pred = predict(mod_1)) %>%
  arrange(pred) %>%
  select(Province_State, pred, deaths_per_thou, cases_per_thou, everything())

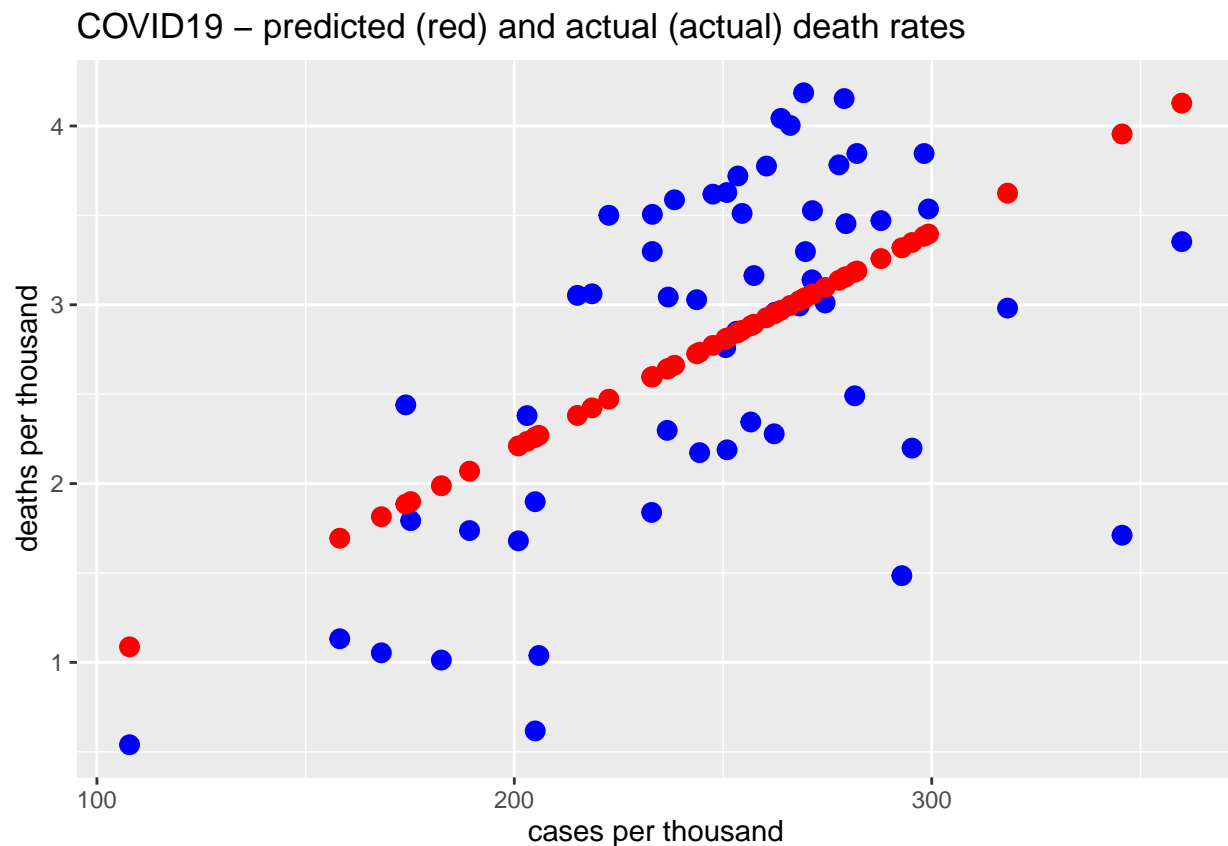
tail(us_state_w_pred, 10)
```

```
## # A tibble: 10 x 7
##   Province_State pred deaths_per_thou cases_per_thou deaths    cases population
##   <chr>          <dbl>          <dbl>          <dbl>    <dbl>    <dbl>    <dbl>
## 1 Wisconsin      3.18            2.49            282.   14502 1639365  5822434
## 2 West Virginia  3.19            3.85            282.    6893  505528  1792147
## 3 South Carolina 3.26            3.47            288.   17869 1481975  5148714
## 4 Utah           3.32            1.49            293.    4761  938864  3205958
## 5 Guam           3.35            2.20            295.     361   48496  164229
```

##	6 Tennessee	3.38	3.85	298.	26265	2036315	6829174
##	7 Kentucky	3.40	3.54	299.	15797	1336858	4467673
##	8 North Dakota	3.62	2.98	318.	2272	242462	762062
##	9 Alaska	3.96	1.71	346.	1268	256089	740995
##	10 Rhode Island	4.13	3.35	360.	3552	381271	1059361

Plot the actual `deaths_per_thou` and the predicted death per thousand (`pred`) and compare the values.

```
us_state_w_pred %>% ggplot +
  geom_point(aes(x = cases_per_thou, y = deaths_per_thou), color = "blue", size = 3) +
  geom_point(aes(x = cases_per_thou, y = pred), color = "red", size = 3) +
  labs(title = "COVID19 - predicted (red) and actual (actual) death rates", x = "cases per thousand", y = "deaths per thousand")
```



The predicted values which were shown in the week 3 lecture that was recorded approximately a year ago demonstrated that `mod_1` did a relatively good job of predicting the `death_per_thou` at the lower and higher ends of the `cases_per_thou` range.

However, this is not necessarily as true when we analyze the data for the last 2 years - probably secondary to a host of additional variables that have been introduced such as more widespread vaccinations and new COVID strains with different degrees of contagiousness and disease severity.

Modeling the Data - part 2

The second linear model will evaluate population density as an independent variable and cases per thousand as the dependent variable using the US data set.

The population densities for the year 2020 are from the “**List of states and territories of the United States by population density**”, link, on Wikipedia.

I could not find this info readily as a csv file and therefore, I chose to add the population density (per kilometer squared) to a vector, `pop_density`, and then add this vector as a column to the `us_states` data set.

```
# Population density of the 56 states (in alphabetical order)

pop_density <- c(38.3, 0.5, 251, 24.3, 22.3, 98.0, 21.5, 288, 196, 4361, 155, 71.9, 283, 87.5, 8.6, 89.

head(pop_density)
```

```
## [1] 38.3 0.5 251.0 24.3 22.3 98.0
```

Add this vector as a column to `us_state_totals`.

```
# arrange the tibble in alphabetical order (bases on Province_State) to match the order in the pop_dens
us_state_totals <- us_state_totals %>%
  arrange(Province_State)

# add the new column of pop_density to the tibble
us_state_totals$Pop_Density <- pop_density

# print out the new tibble and check that the correct values in the Pop_Density column have been associ
us_state_totals %>%
  select(Province_State, Pop_Density, everything())
```

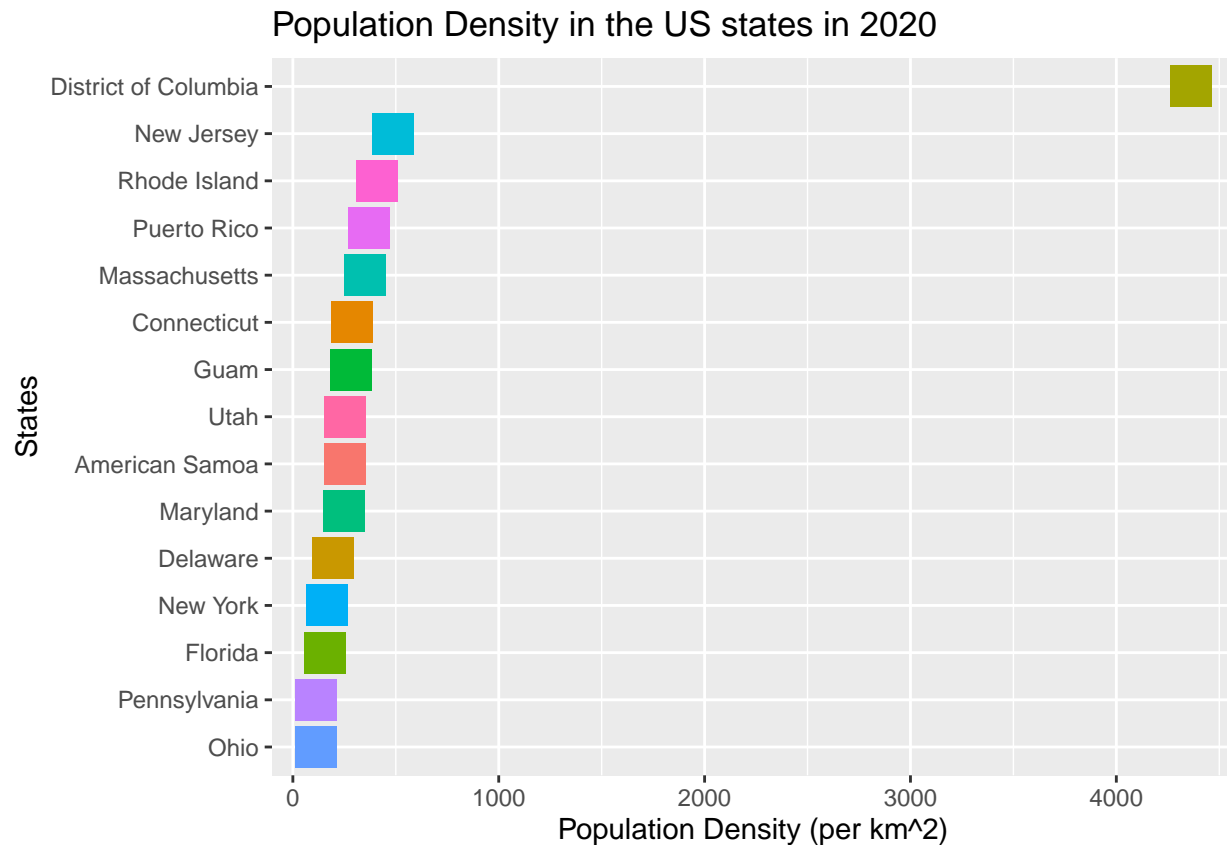
```
## # A tibble: 56 x 7
##   Province_State      Pop_Density deaths    cases population cases_per_thou
##   <chr>            <dbl>   <dbl>   <dbl>      <dbl>         <dbl>
## 1 Alabama          38.3  19628 1304710   4903185         266.
## 2 Alaska           0.5   1268  256089    740995         346.
## 3 American Samoa   251     30    5999    55641         108.
## 4 Arizona          24.3  30230 2030925   7278717         279.
## 5 Arkansas          22.3  11415  838251   3017804         278.
## 6 California        98   90782 9349673  39512223         237.
## 7 Colorado          21.5  12513 1407405   5758736         244.
## 8 Connecticut       288   10914  779460   3565287         219.
## 9 Delaware          196    2931  267265    973764         274.
## 10 District of Columbia 4361    1340  144675    705749         205.
## # ... with 46 more rows, and 1 more variable: deaths_per_thou <dbl>
```

Observation: The District Of Columbia is somewhat of an outlier with a population density of 4361 per kilometers squared. New Jersey has the second highest population density.

Evaluate the range of values for population density in the 56.

```
us_state_totals %>%
  filter(Pop_Density > 100) %>%
  ggplot(aes(x = fct_reorder(Province_State, Pop_Density), Pop_Density, group = Province_State, color =
  geom_point(show.legend = FALSE, size = 7, shape = 15) +
```

```
coord_flip() +
theme(legend.position="bottom",
      axis.text.x = element_text(angle = 0)) +
labs(title = "Population Density in the US states in 2020", x = "States", y = "Population Density (per
```



Observation: the population density for District of Columbia (DC) is an outlier.

Model 2 Pop_Density as a predictor for cases_per_thou and deaths_per_thou (a linear model)

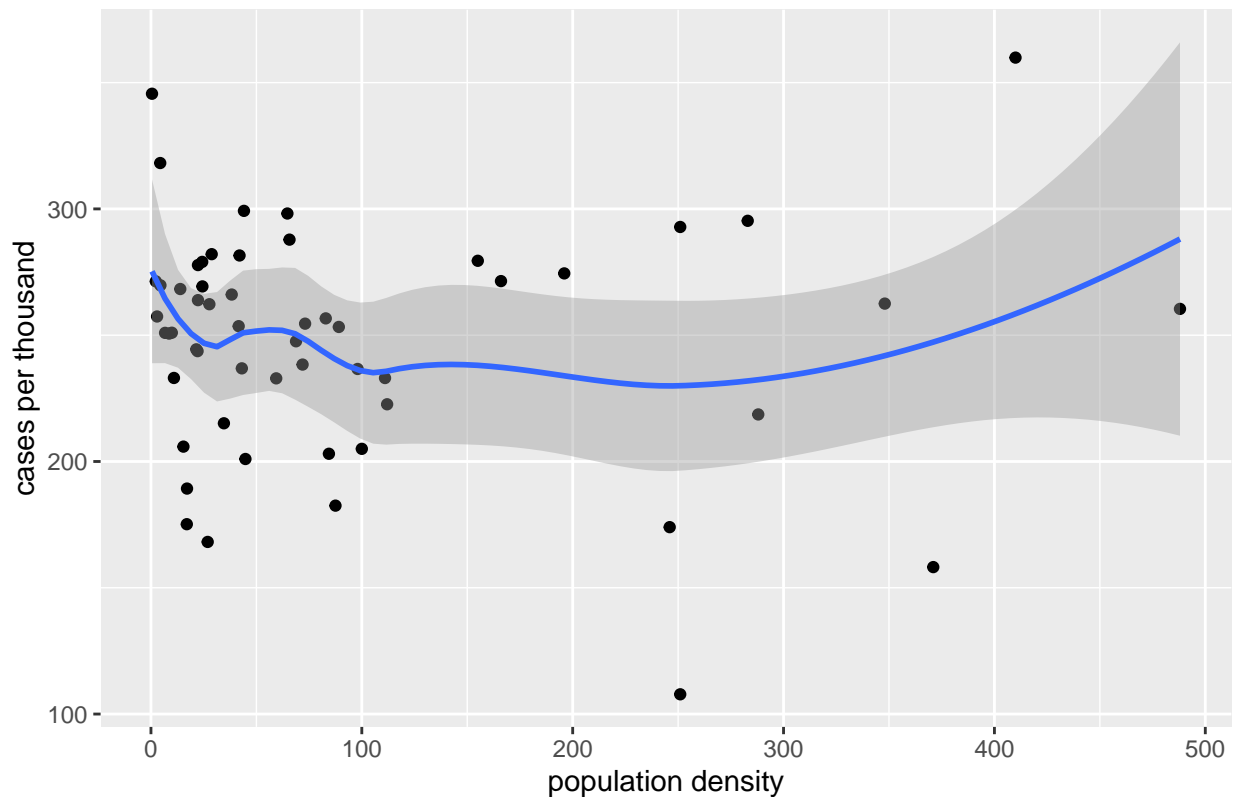
Plot the two variables as a scatter plot to see their relationship. As the population density for DC is an outlier, I will not include DC in the model.

```
us_state_totals_minusDC <- us_state_totals %>%
  filter(Province_State != "District of Columbia")

us_state_totals_minusDC %>%
  ggplot(aes(Pop_Density, cases_per_thou)) +
  geom_point() +
  geom_smooth() +
  labs(title = "COVID19 - population density as a predictor of cases rates per thousand", x = "population density", y = "cases rates per thousand")

## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```

COVID19 – population density as a predictor of cases rates per thousand



Linear model for the two variables.

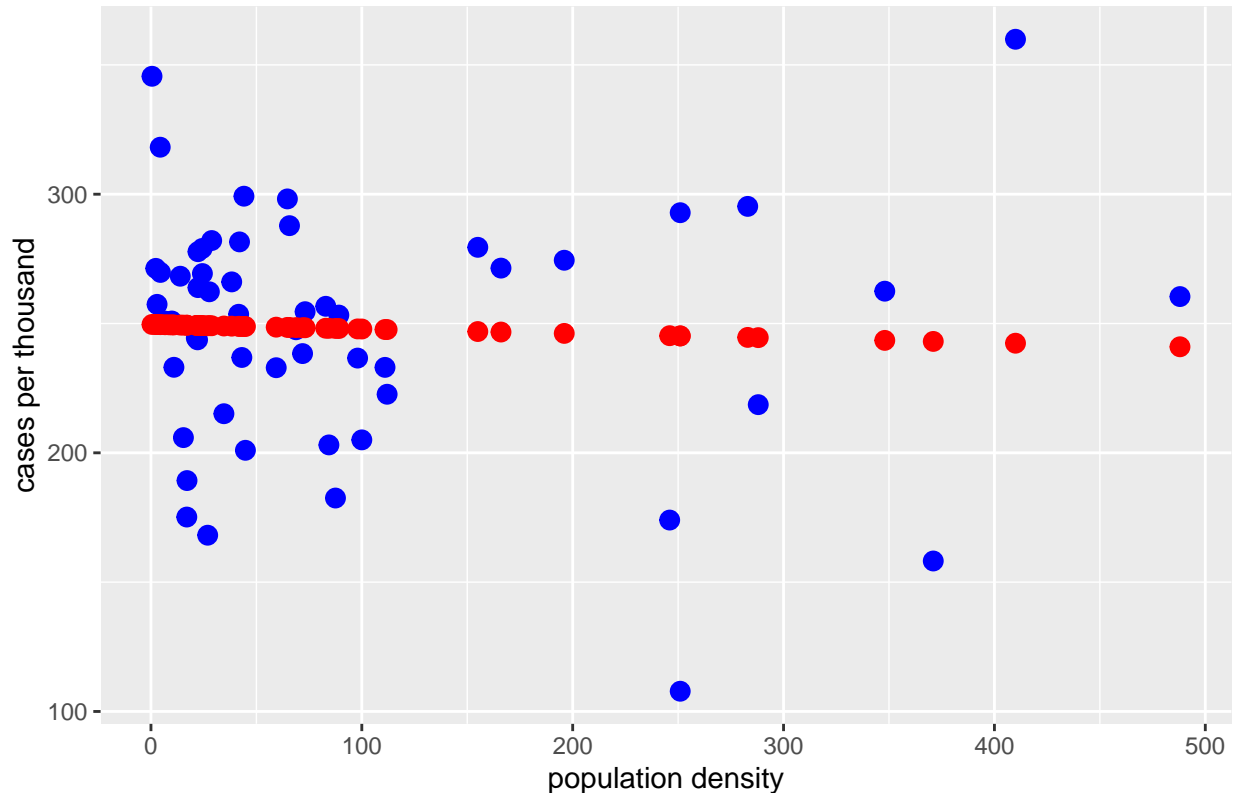
```
mod_2 <- lm(cases_per_thou ~ Pop_Density, data = us_state_totals_minusDC)
summary(mod_2)
```

```
##
## Call:
## lm(formula = cases_per_thou ~ Pop_Density, data = us_state_totals_minusDC)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -137.383  -20.698    5.198   26.504  117.524
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  249.64726    8.04111   31.046  <2e-16 ***
## Pop_Density  -0.01772    0.05389   -0.329    0.744
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 45.87 on 53 degrees of freedom
## Multiple R-squared:  0.002036, Adjusted R-squared:  -0.01679
## F-statistic: 0.1081 on 1 and 53 DF, p-value: 0.7436
```

Plot the prediction for cases_per_thou based on Pop_Density.

```
us_state_w_pred_2 <- us_state_totals_minusDC %>%
  mutate(cases_pred = predict(mod_2)) %>%
  ggplot() +
    geom_point(aes(x = Pop_Density, y = cases_per_thou), color = "blue", size = 3) +
    geom_point(aes(x = Pop_Density, y = cases_pred), color = "red", size = 3) +
    labs(title = "COVID19 - predicted (red) and actual (actual) case rates", x = "population density", y = "cases per thousand")
us_state_w_pred_2
```

COVID19 – predicted (red) and actual (actual) case rates



Observation: population density does not explain the variations in cases per thousand across different states. Note that this is overall a very big picture view of the problem as more optimally population density in different counties, cities or even neighborhoods could be used as a predictive variable (in conjunction with other predictive variables such as household income, ...).

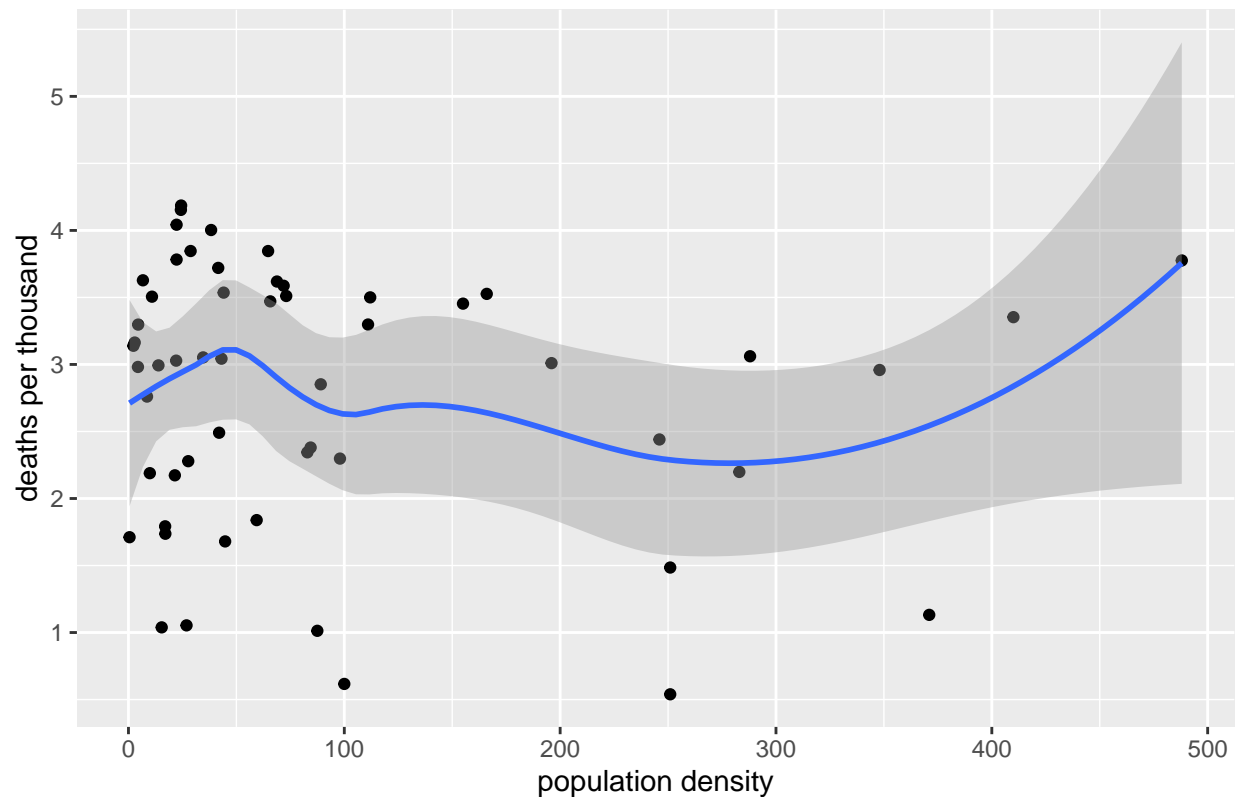
Model 3 Pop_Density as a predictor for `deaths_per_thou` (a linear model)

The same process can be used to evaluate the relationship between these two variables. I will replicate the steps in Model 2 as a chain of code chunks for this model.

```
us_state_totals_minusDC %>%
  ggplot(aes(Pop_Density, deaths_per_thou)) +
    geom_point() +
    geom_smooth() +
    labs(title = "COVID19 - population density as a predictor of death rates per thousand", x = "population density", y = "deaths per thousand")
```

```
## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```

COVID19 – population density as a predictor of death rates per thousand

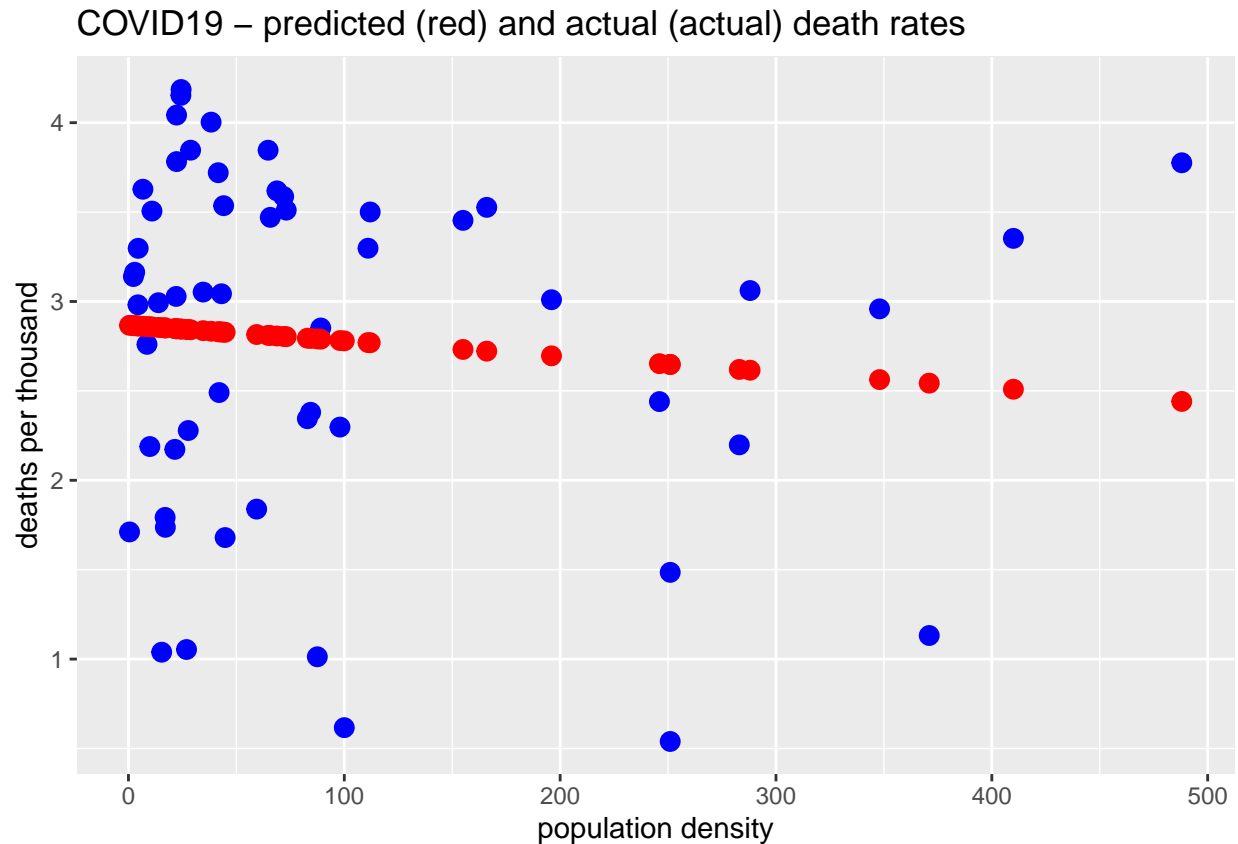


```
mod_3 <- lm(deaths_per_thou ~ Pop_Density, data = us_state_totals_minusDC)
summary(mod_3)
```

```
##
## Call:
## lm(formula = deaths_per_thou ~ Pop_Density, data = us_state_totals_minusDC)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.1632 -0.6171  0.2157  0.7489  1.3398
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.867110   0.169394  16.926  <2e-16 ***
## Pop_Density -0.000873   0.001135  -0.769    0.445
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9663 on 53 degrees of freedom
## Multiple R-squared:  0.01104,    Adjusted R-squared:  -0.007623
## F-statistic: 0.5915 on 1 and 53 DF,  p-value: 0.4453
```

```
us_state_w_pred_3 <- us_state_totals_minusDC %>%
  mutate(deaths_pred = predict(mod_3)) %>%
```

```
ggplot() +
  geom_point(aes(x = Pop_Density, y = deaths_per_thou), color = "blue", size = 3) +
  geom_point(aes(x = Pop_Density, y = deaths_pred), color = "red", size = 3) +
  labs(title = "COVID19 - predicted (red) and actual (actual) death rates", x = "population density", y =
us_state_w_pred_3
```



Observation: Similar to Model 2, population density does not explain the variations in deaths per thousand across different states.

Sources of Bias

The sources of bias in the US and global data sets are due to how the data were collected and reported and a multitude of other factors. Different states and countries may have different criteria for counting an individual as having contracted COVID or attributing a death to COVID. The population density and lock down and masking policies effect the transmission of the virus in the community as well the cohorts who are more likely to get exposed. Another source of bias is the different strains of the COVID virus - these may all not have the same degree of contagiousness or cause the comparable disease severity. Access to the COVID vaccine and the timing of vaccination during the pandemic have varied both between and within countries - moreover, political affiliations have clouded individual's choices in receiving the vaccine.

The reporting bias could result in under reporting of both cases and deaths. These data also do not capture the age of those affected and those who succumbed to the disease - the susceptibility of different age groups to the evolving strains may also have changed over time.

Conclusion

For this assignment I replicated the code that used by Dr.Wall in the week three lectures and added my own analyses, visualizations and models. The assignment is somewhat lengthy as I have tried to meticulously document the exploratory and analytic steps. Moreover, I am new to R and I have learned a lot from the impressive online resources available for R programming and tidyverse which I tried to implement in the analyses and visualizations.

This assignment reinforced how complex data wrangling, analysis, visualization and modeling can be. I look forward to learning more on these subjects in future courses.

```
sessionInfo()
```

Session Info

```
## R version 4.1.2 (2021-11-01)
## Platform: x86_64-w64-mingw32/x64 (64-bit)
## Running under: Windows 10 x64 (build 22000)
##
## Matrix products: default
##
## locale:
## [1] LC_COLLATE=English_United States.1252
## [2] LC_CTYPE=English_United States.1252
## [3] LC_MONETARY=English_United States.1252
## [4] LC_NUMERIC=C
## [5] LC_TIME=English_United States.1252
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods    base
##
## other attached packages:
## [1] lubridate_1.8.0 forcats_0.5.1  stringr_1.4.0  dplyr_1.0.7
## [5] purrr_0.3.4     readr_2.1.1    tidyr_1.1.4    tibble_3.1.6
## [9] ggplot2_3.3.5   tidyverse_1.3.1
##
## loaded via a namespace (and not attached):
## [1] Rcpp_1.0.8      lattice_0.20-45 assertthat_0.2.1 digest_0.6.29
## [5] utf8_1.2.2      R6_2.5.1        cellranger_1.1.0 backports_1.4.1
## [9] reprex_2.0.1    evaluate_0.14   httr_1.4.2      highr_0.9
## [13] pillar_1.6.4    rlang_1.0.2     curl_4.3.2      readxl_1.3.1
## [17] rstudioapi_0.13 Matrix_1.3-4     rmarkdown_2.14  splines_4.1.2
## [21] labeling_0.4.2  bit_4.0.4       munsell_0.5.0   broom_0.7.11
## [25] compiler_4.1.2  modelr_0.1.8    xfun_0.30       pkgconfig_2.0.3
## [29] mgcv_1.8-38     htmltools_0.5.2 tidyselect_1.1.1 fansi_1.0.2
## [33] crayon_1.4.2    tzdb_0.2.0      dbplyr_2.1.1    withr_2.4.3
## [37] grid_4.1.2      nlme_3.1-157    jsonlite_1.7.3  gtable_0.3.0
## [41] lifecycle_1.0.1 DBI_1.1.2       magrittr_2.0.1  scales_1.2.0
## [45] cli_3.1.0        stringi_1.7.6   vroom_1.5.7     farver_2.1.0
## [49] fs_1.5.2         xml2_1.3.3      ellipsis_0.3.2  generics_0.1.1
## [53] vctrs_0.3.8      tools_4.1.2     bit64_4.0.5     glue_1.6.0
```

```
## [57] hms_1.1.1      parallel_4.1.2  fastmap_1.1.0   yaml_2.2.1
## [61] colorspace_2.0-2 rvest_1.0.2     knitr_1.37      haven_2.4.3
```